



Trustworthy Machine Learning on Imbalance Data



Jiangchao Yao

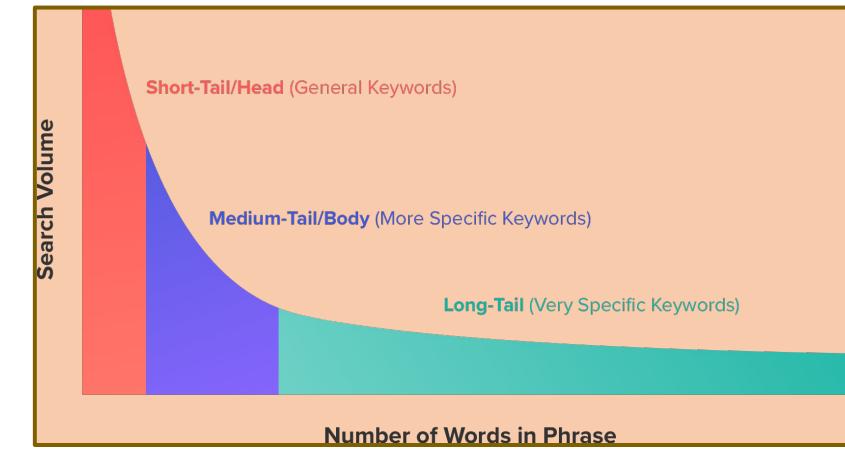
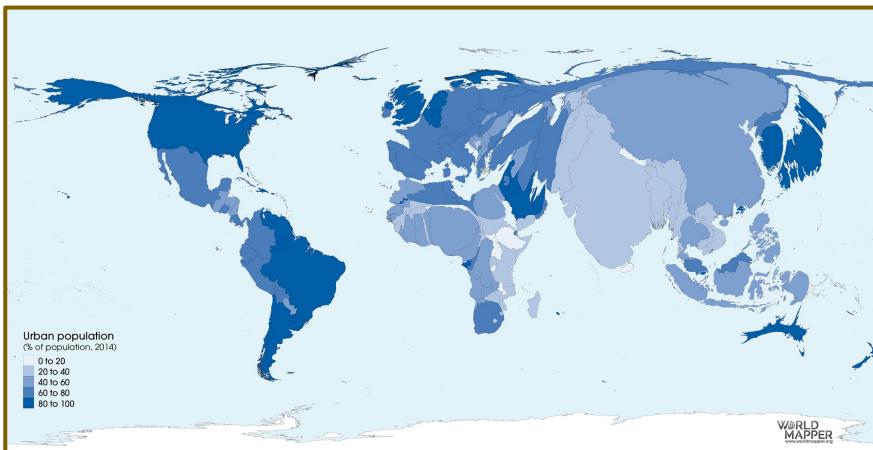
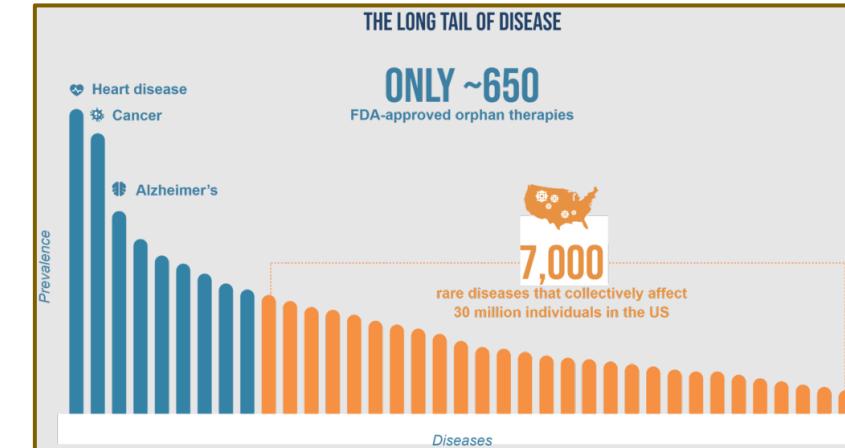
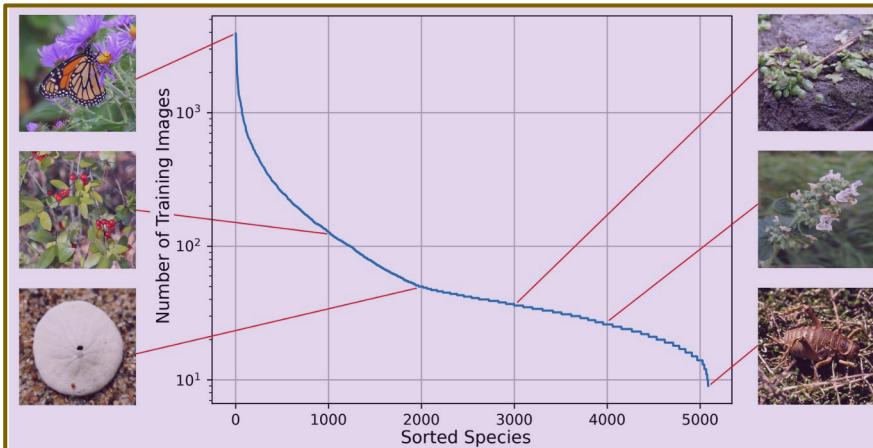
CMIC, Shanghai Jiao Tong University



The Imbalance Nature of Data



Large-scale natural sources are very imbalance, usually following a **long-tailed distribution**.



[1] Van Horn et al. The iNaturalist Species Classification and Detection Dataset. CVPR 2018.

[3] <https://worldmapper.org/maps/urban-population-relative-2014/>

[2] Gregory et al. CXR-LT challenge. ICCV CVAMD 2023.

[4] <https://seopressor.com/blog/short-tail-or-long-tail-keywords/>

How Imbalance Data Matters in Machine Learning



A direct decomposition on the risk minimization

$$\min_f R(f) = E_{P(x,y)}[\ell(f(x), y)] = \sum_{k=1}^K P(y=k) E_{P(x|y=k)}[\ell(f(x), y)]$$

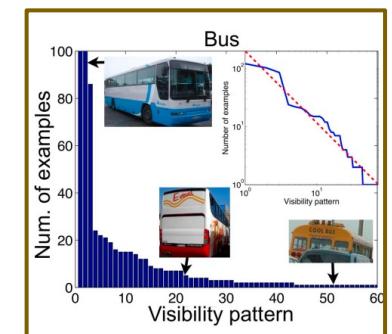
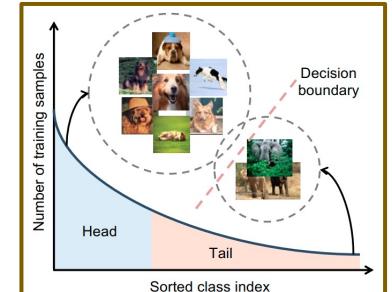


Minority (“generalized” conceptual) classes have weak importance for training, which will be easily ignored in the early phase especially for overparameterized DNNs [1] (*or namely, will be sacrificed first if it is not sufficient for the model to learn*).



However, in real applications, the value of classes cannot be absolutely characterized by their quantity, and instead, sometimes **less is more** for sustainable long-term development.

- **Fairness** w.r.t. diversity e.g., small populations of gender, race and consumers
- **Cost-sensitive scenarios** e.g., medical disease diagnosis and treatment (COVID-19)



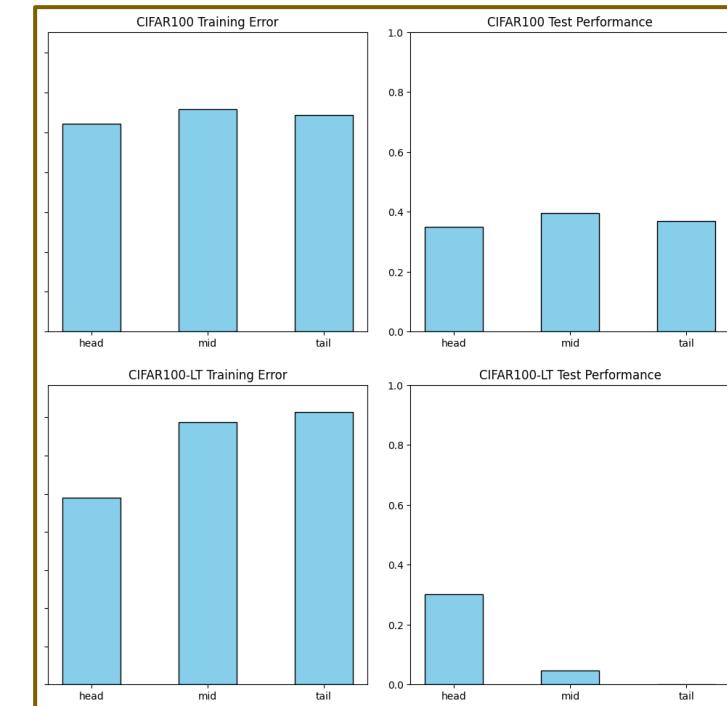
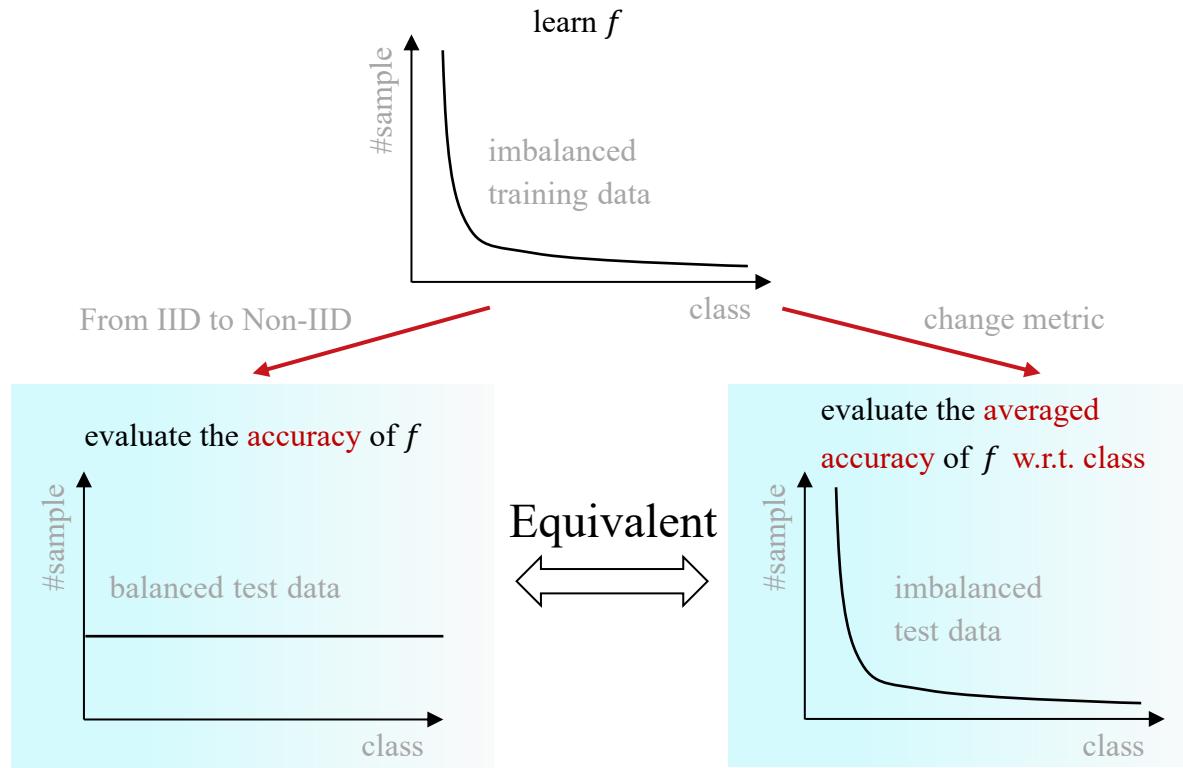
[1] Vitaly Feldman. "Does Learning Require Memorization? A Short Tale about a Long Tail." SIGACT 2020.



Why the Resulted Imbalanced Learning is Special



A critical highlight on the evaluation, different from the ordinary IID learning



The change in evaluation metric induces **an statistical consistency problem** on applying conventional learning methods, that is,

What we design during training should be statistically consistent with what we pursue about the evaluation.





The Historical Development

The development of imbalance learning

1998

2009

2018

2023

Optimizing Classifiers for Imbalanced Training Sets

Grigoris Karakoulas
Oxford Internet Institute,
Canadian Imperial Bank of Commerce
161 Bay St., B2C-111
Toronto, ON, Canada M5A 2S8
Email: gkarakou@cs.ox.ac.uk

John Shawe-Taylor
Department of Computer Science,
Royal Holloway, University of London
Egham, TW20 0EX
Email: jst@cs.rhul.ac.uk

Abstract
Following recent results [9, 8] showing the importance of the fat-shattering dimension in explaining the beneficial effect of a large margin on generalization performance, the present paper investigates the use of this measure to help set the threshold for separating points in a minority class from the majority class. The proposed method is a relatively new challenge that has attracted growing attention from both academics and industry. The imbalanced learning problem is a well-known challenge in machine learning due to the inherent class distribution skewness. Due to the inherent class distribution skewness, learning from imbalanced data sets is a challenging task. In this paper, we provide a comprehensive review of the development of research in learning from imbalanced data sets. We also discuss the challenges of learning from imbalanced data sets and propose some potential solutions. Finally, we introduce some metrics used to evaluate learning performance under the imbalanced learning paradigm. Furthermore, in order to evaluate learning performance under the imbalanced learning paradigm, we propose some evaluation metrics and challenges, as well as potential research directions for learning from imbalanced data.

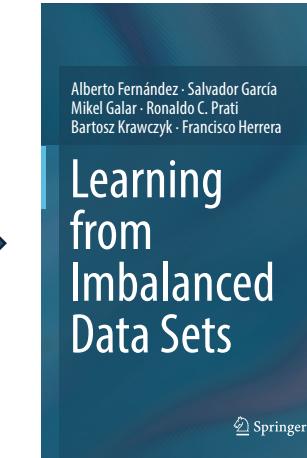
methods

Learning from Imbalanced Data

Habib He, Member, IEEE, and Edwardo A. Garcia

Abstract—With the continuous expansion of data availability in many large-scale, complex, and imbalanced systems, such as surveillance, security, Internet, and finance, it becomes critical to advance the fundamental understanding of knowledge discovery and data mining in these domains. This paper reviews the state-of-the-art research in learning from imbalanced data sets. The imbalanced learning problem is a well-known challenge in machine learning due to the inherent class distribution skewness. Due to the inherent class distribution skewness, learning from imbalanced data sets is a challenging task. In this paper, we provide a comprehensive review of the development of research in learning from imbalanced data sets. We also discuss the challenges of learning from imbalanced data sets and propose some potential solutions. Finally, we introduce some metrics used to evaluate learning performance under the imbalanced learning paradigm. Furthermore, in order to evaluate learning performance under the imbalanced learning paradigm, we propose some evaluation metrics and challenges, as well as potential research directions for learning from imbalanced data.

Applications

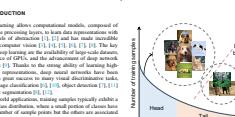


Deep learning

Deep Long-Tailed Learning: A Survey

Yiyan Zhang, Binji Kang, Bryan Hooi, Shucheng Yan, Fellow, IEEE, and Jiashi Feng

Abstract—Deep long-tailed learning is one of the most challenging problems in visual recognition, due to its well-performing deep learning models that are trained on large-scale datasets with balanced classes. However, deep learning models trained on balanced datasets are often incapable of learning high-quality image representations and fail to generalize well enough to handle the long-tailed learning problem. To address this issue, this paper provides a comprehensive survey of deep learning models for real-world applications, where they can be easily trained to learn distinct classes and perform well on long-tailed learning tasks. Specifically, this paper first introduces the basic concepts of long-tailed learning and reviews the main research areas in long-tailed learning. To be specific, we group existing deep long-tailed learning models into three main categories: (1) deep learning models for learning from long-tailed data; (2) deep learning models for learning from short-tailed data; and (3) deep learning models for learning from balanced data. Additionally, we empirically analyze several state-of-the-art methods by evaluating to what extent they address the issue of long-tailed learning. Finally, we provide some future research directions for learning from long-tailed learning and identifying several promising directions for future research.

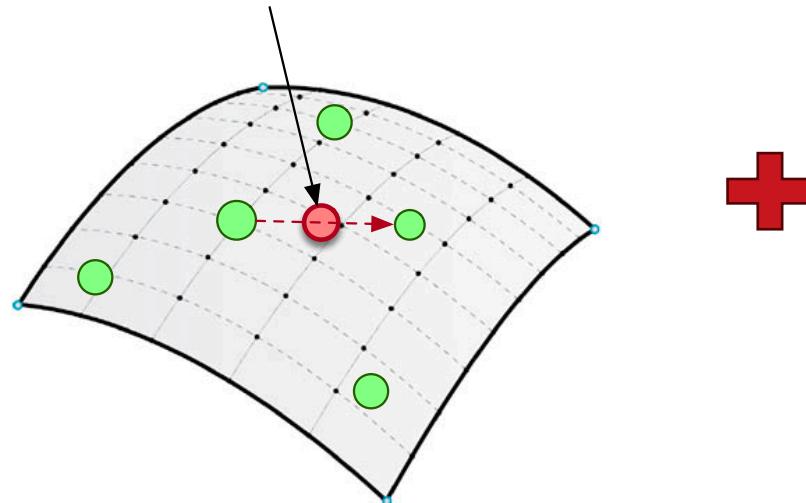


- [1] Karakoulas et al. Optimizing Classifiers for Imbalanced Training Sets. NIPS 1998.
- [2] He et al. Learning from Imbalanced Data. TKDE 2009.
- [3] Fernández et al. Learning from Imbalanced Data Sets. Springer, 2018.
- [4] Zhang et al. Deep Long-tailed Learning: A Survey. TPAMI 2023.

SMOTE: Synthetic Minority Over-sampling Technique

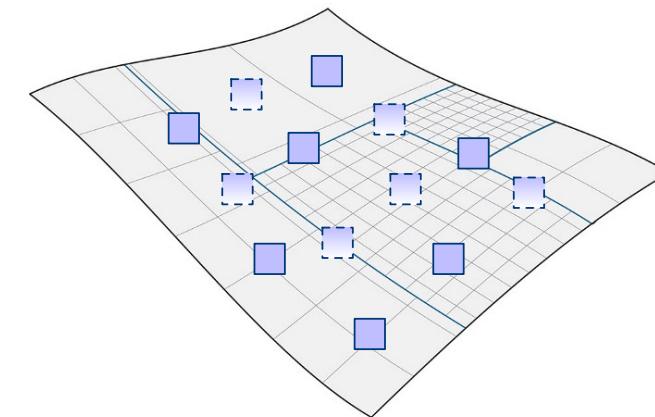
Motivation: Replication of the minority class does not cause its decision boundary to spread into the majority class region (but overfitting).

Interpolation on minority manifold



The main idea of SMOTE: augmentation for minority class by interpolation instead of over-sampling with replacement.

Under-sampling in the majority class



Interpolation is limited by the samples. Thus, SMOTE also always runs with the under-sampling for majority class.

Threshold-Moving: adjust the prediction in a post-hoc manner.

Motivation: The over-confident prediction for majority or the low-confident prediction for minority can be calibrated after training.

THE THRESHOLD-MOVING ALGORITHM

Training phase:

1. Let S be the original training set.
2. Train a neural network from S .

Test phase:

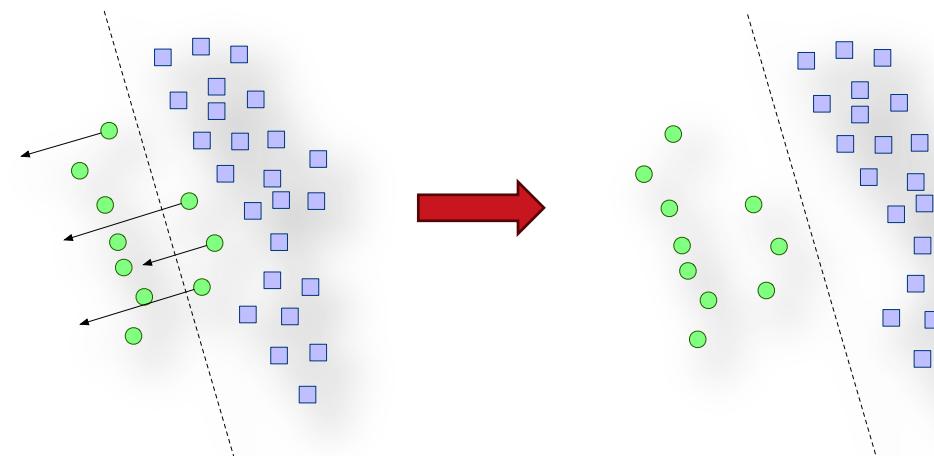
1. Generate real-value outputs with the trained neural network.
2. For every output, multiply it with the sum of the costs of misclassifying the corresponding class to other classes.
3. Return the class with the biggest output.

Moving function

$$\hat{p}_k = \frac{p_k * \sum_{k'=1}^K C[k][k']}{\eta}$$

where p_k is the probabilistic prediction, $C[k][k']$ is the cost mis-predicted from class k to k' , and η is renormalization parameter.

Sampling methods might not always show promise in multi-class imbalance learning, but threshold-moving way does.



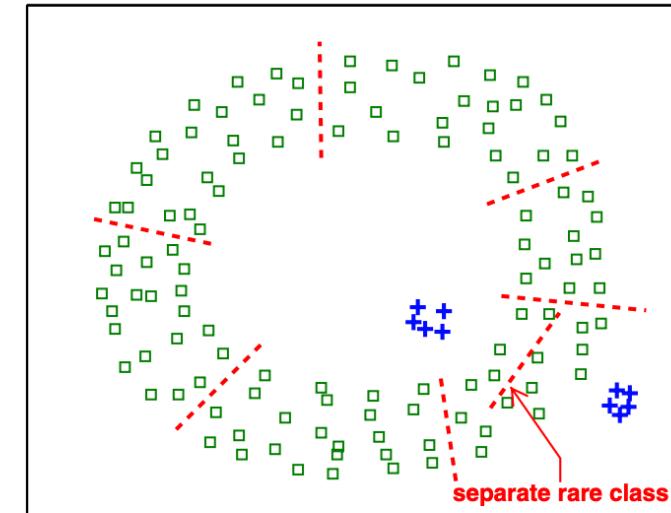
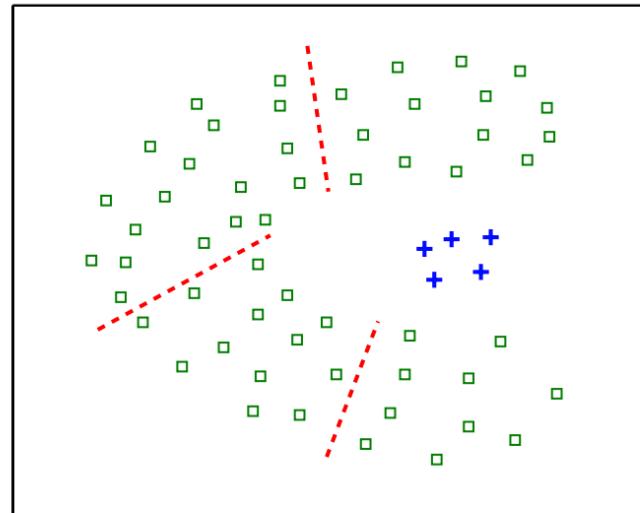
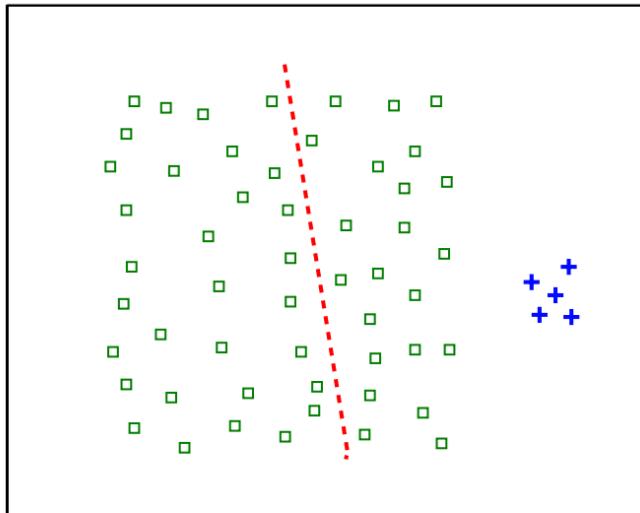
Let $p_k = \frac{e^{z_k}}{\sum_{k'} e^{z_{k'}}}$ denote the class prediction. If we set $\sum_{k'=1}^K C[k][k'] = e^{-\tau \log \pi_k}$ where π_k is the class prior and τ is the temperature, the threshold-moving method recovers the popular **logit adjustment** method for long-tailed learning.

Majority classes have the smaller cost than minority classes, e.g., $e^{-\tau \log \pi_k}$ is monotonously decreasing.

COG: Local Decomposition for Rare Class Analysis

Intuition: Quantity imbalance limits the learning pace of minority over majority. We can adjust the quantities by decomposition.

How to properly decompose the majority classes (or including minority classes) into subclasses to balance the training?



Phase I: local clustering

1. for class $i = 1$ to c // “ c ” represents #classes
2. $\text{clusterLabel}(i) = \text{Clustering}(\mathcal{D}(i), \mathbf{K}(i));$
3. $\mathcal{D}(i)^* = \text{changeLabel}(\mathcal{D}(i), \text{clusterLabel}(i));$
4. end for

Phase II: over-sampling (for COG-OS only)

5. for class $j = 1$ to c
6. $\mathcal{D}(j)^{**} = \text{replicate}(\mathcal{D}(j)^*, r(j))$
7. end for
8. $\mathcal{D}^{**} = \bigcup_{j=1}^c (\mathcal{D}(j)^{**});$

Phase III: training

9. $\mathbb{M} = \text{train}(\mathcal{D}^{**}, \mathbb{L});$

Phase IV: predicting

10. $\mathbf{p}' = \text{predict}(\mathcal{T}, \mathbb{M});$
11. $\mathbf{p} = \text{convertLabel}(\mathbf{p}');$



Retrospection-IV: Theory for Imbalance Learning



On Statistical Consistency of Binary Classification with Balanced Accuracy

Motivation: The early ERM theory is developed for the instance-wise evaluation, but cannot guarantee the consistency for balanced measure.

$$\text{Accuracy} = \mathbb{E}_{p(x,y)}[h(x) = y]$$



$$\text{Balanced Accuracy} = \frac{\sum_{k \in \{-1,1\}} \mathbb{E}_{p(x|y=k)}[h(x)=k]}{2}$$

If we consider the balanced accuracy, how to modify the algorithm to satisfy the statistical consistency?

Theorem 3. Let D be a probability distribution on $\mathcal{X} \times \{\pm 1\}$ satisfying Assumption A. Let \hat{p}_S denote any estimator of $p = \mathbf{P}(y = 1)$ satisfying $\hat{p}_S \in (0, 1)$ and $\hat{p}_S \xrightarrow{P} p$. Let $\hat{\eta}_S : \mathcal{X} \rightarrow [0, 1]$ denote any class probability estimator satisfying $\mathbf{E}_x[|\hat{\eta}_S(x) - \eta(x)|^r] \xrightarrow{P} 0$ for some $r \geq 1$, and let $h_S(x) = \text{sign}(\hat{\eta}_S(x) - \hat{p}_S)$. Then

$$\text{regret}_D^{\text{AM}}[h_S] \xrightarrow{P} 0.$$

Algorithm 1 Plug-in with Empirical Threshold

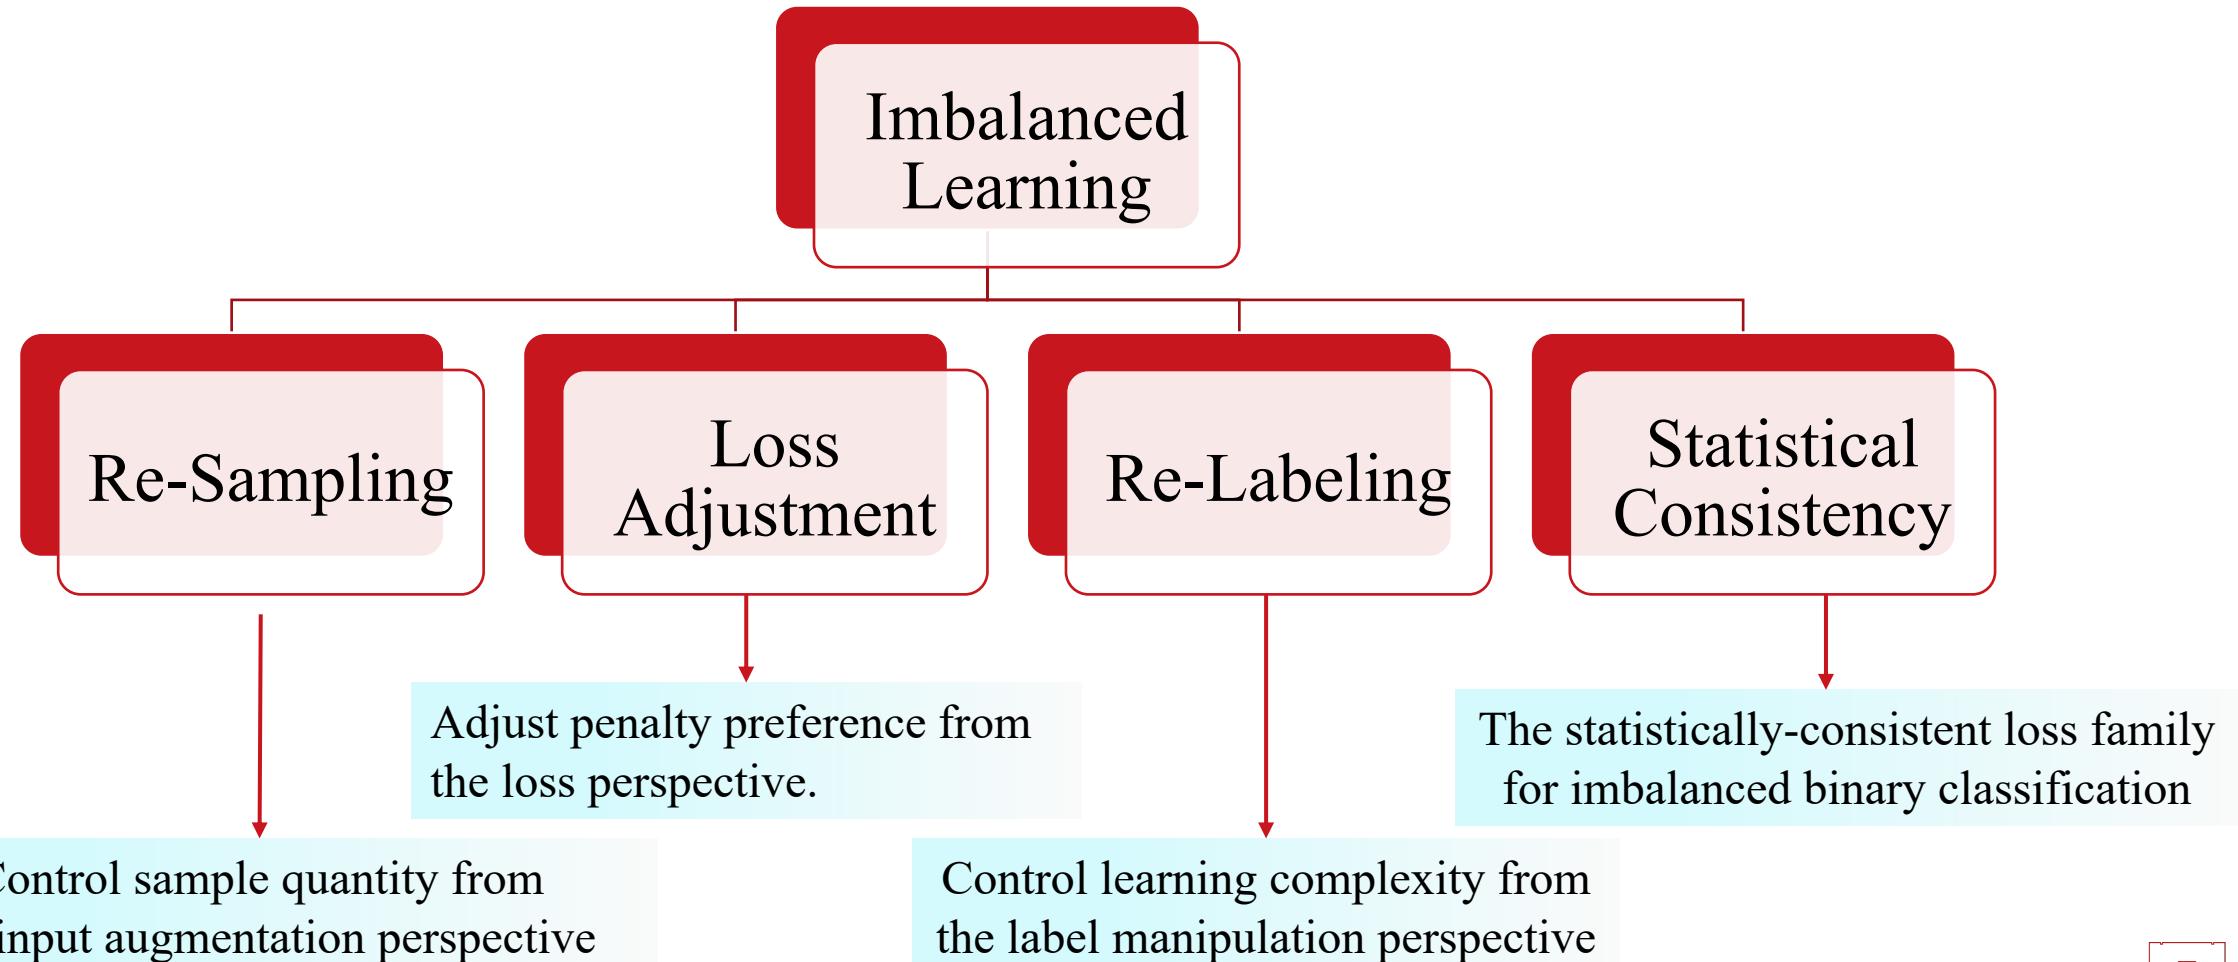
- 1: **Input:** $S = ((x_1, y_1), \dots, (x_n, y_n)) \in (\mathcal{X} \times \{\pm 1\})^n$
- 2: **Select:** (a) Proper (composite) loss $\ell : \{\pm 1\} \times \bar{\mathbb{R}} \rightarrow \bar{\mathbb{R}}_+$, with link function $\psi : [0, 1] \rightarrow \bar{\mathbb{R}}$; (b) RKHS \mathcal{F}_K with positive definite kernel $K : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}$; (c) regularization parameter $\lambda_n > 0$
- 3: $f_S \in \operatorname{argmin}_{f \in \mathcal{F}_K} \left\{ \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i)) + \lambda_n \|f\|_K^2 \right\}$
- 4: $\hat{\eta}_S = \psi^{-1} \circ f_S$
- 5: $\hat{p}_S = \text{(as in Eq. (2))}$
- 6: **Output:** Classifier $h_S(x) = \text{sign}(\hat{\eta}_S(x) - \hat{p}_S)$



Summary



Summary of imbalanced learning in the early years





What is the new of this topic in the recent years?

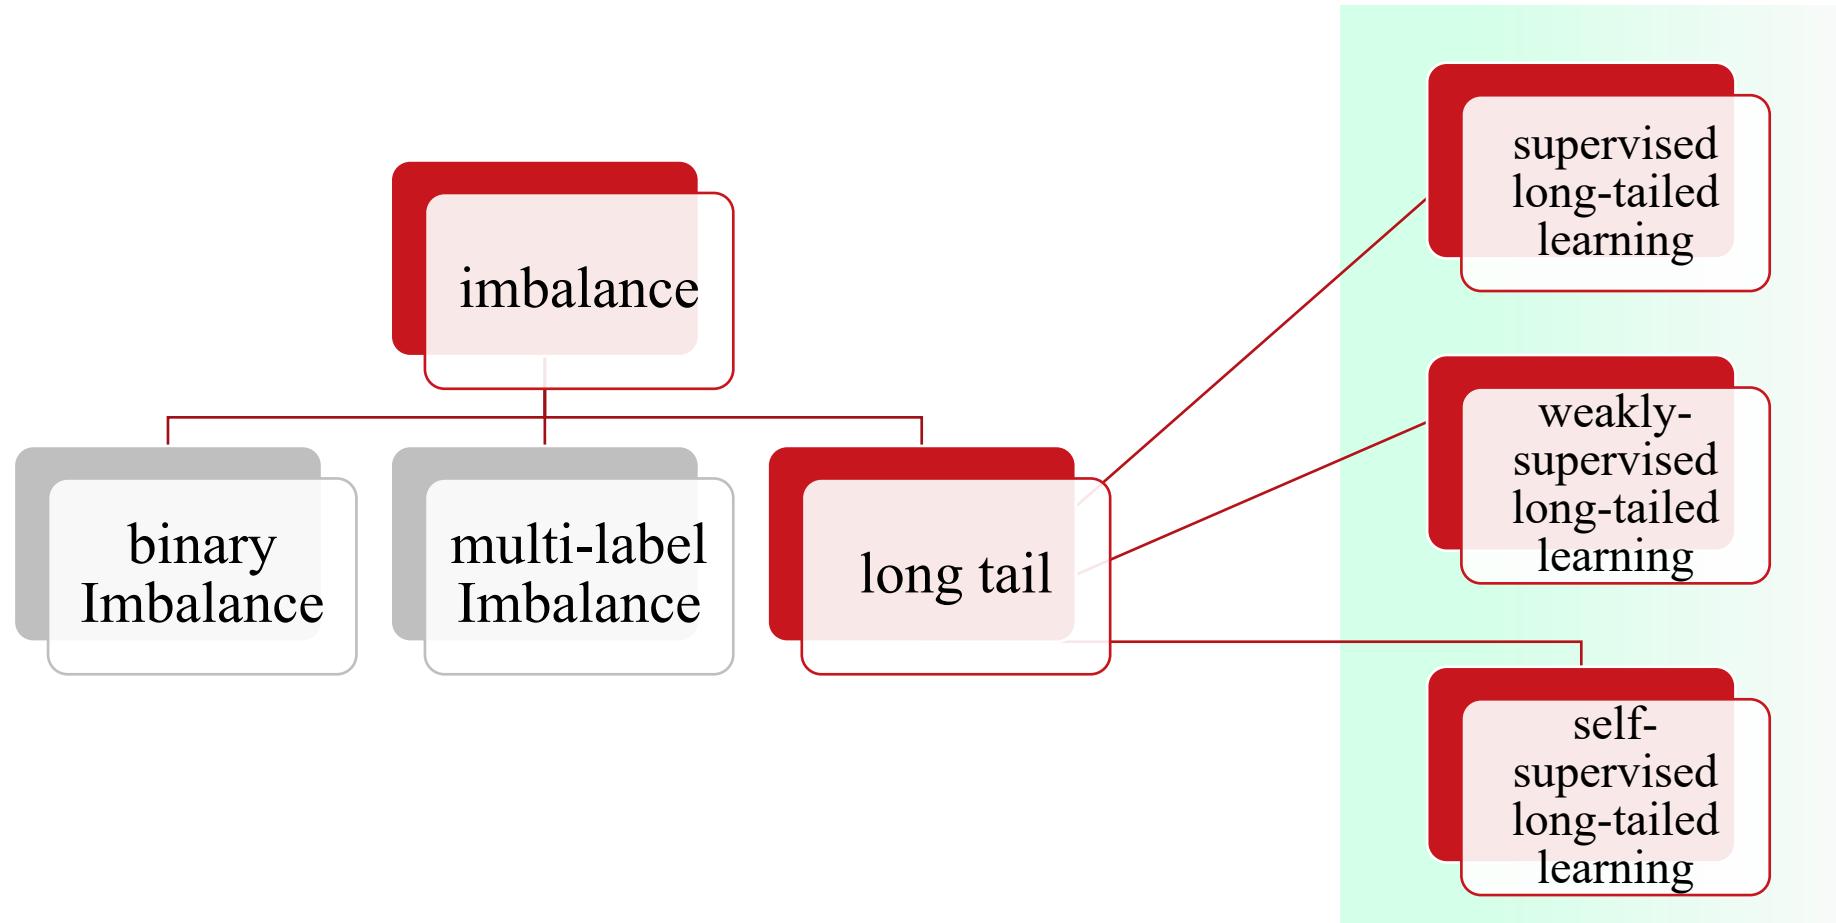




The Following Part in This Tutorial



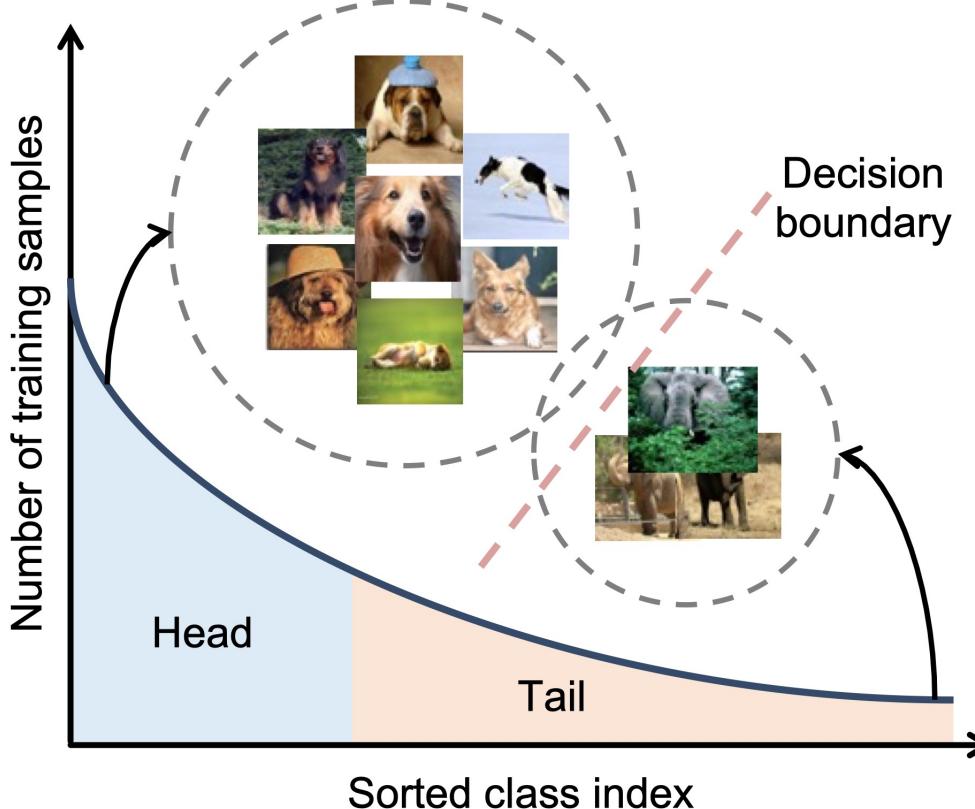
The recent advances of imbalance learning powered by deep learning



Supervised Long-tailed Learning



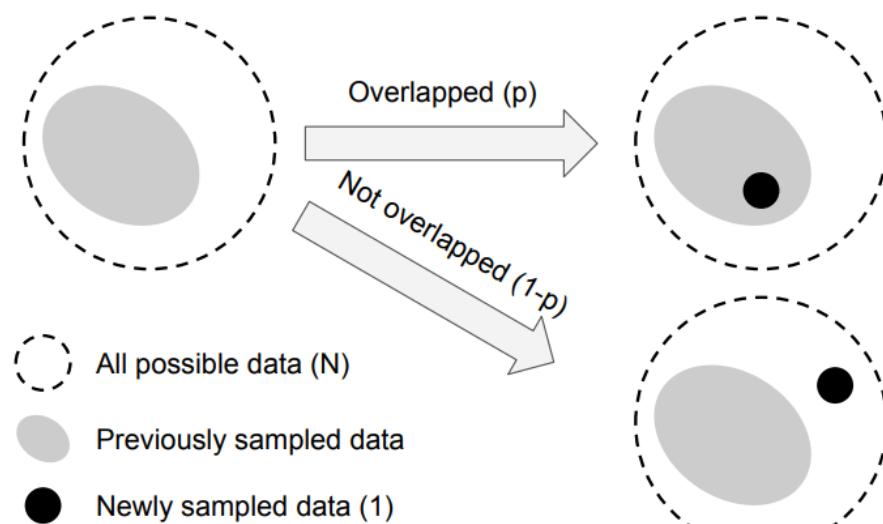
**It has been contributed with
very broad explorations**



Method	Year	Class Re-balancing			Augmentation		Module Improvement			Target Aspect
		Re-sampling	CSL	LA	TL	Aug	RL	CD	DT	
LMLE [89]	2016						✓			feature
HFL [90]	2016						✓			feature
Focal loss [54]	2017				✓					objective
Range loss [21]	2017									feature
CRL [50]	2017							✓		feature
MetaModelNet [91]	2017									
DSTL [92]	2018						✓			sample
DCL [93]	2019							✓		objective
Meta-Weight-Net [94]	2019						✓			objective
LDAM [18]	2019							✓		objective
CB [16]	2019						✓			feature
UML [95]	2019						✓			feature
FTL [96]	2019							✓		feature
Unequal-training [48]	2019							✓		feature
OLTR [15]	2019							✓		feature
Balanced Meta-Softmax [97]	2020							✓		feature
Decoupling [32]	2020							✓		feature
LST [98]	2020							✓		feature
Domain adaptation [28]	2020							✓		feature
Equalization loss (ESQL) [19]	2020							✓		feature
DBM [22]	2020							✓		sample
Distribution-balanced loss [37]	2020							✓		objective
UNO-IC [99]	2020							✓		objective
De-confound-TDE [45]	2020							✓		prediction
M2m [100]	2020							✓		sample
LEAP [49]	2020							✓		feature
OFA [101]	2020							✓		feature
SSP [102]	2020							✓		sample, model
LFME [103]	2020							✓		feature
IEM [104]	2020							✓		classifier
Deep-RTC [105]	2020							✓		sample, model
SimCal [34]	2020							✓		sample, model
BBN [44]	2020							✓		sample, model
BAGS [56]	2020							✓		sample
VideoLT [38]	2021						✓			sample, objective
LOCE [33]	2021						✓			sample
DARS [26]	2021						✓			classifier
CREST [106]	2021						✓			feature
GIST [107]	2021						✓			objective
FASA [58]	2021						✓			objective
Equalization loss v2 [108]	2021						✓			sample
Sesasaw loss [109]	2021						✓			objective
ACSL [110]	2021						✓			objective
IB [111]	2021						✓			objective
PML [51]	2021						✓			objective
VS [112]	2021						✓			objective
LADE [31]	2021						✓			objective
RobBal [113]	2021						✓			objective
DisAlign [29]	2021						✓			classifier
MiLAS [114]	2021						✓			objective, feature, classifier
Logit adjustment [14]	2021						✓			prediction
Conceptual 12M [115]	2021									feature
DIVE [116]	2021									model
MosaicOS [117]	2021									sample
RSG [118]	2021									feature
SSD [119]	2021									feature
RIDE [17]	2021									feature
MetaAug [120]	2021									feature
PaCo [121]	2021									classifier
DRO-LT [122]	2021									sample, model
Unsupervised discovery [35]	2021									sample, model
Hybrid [123]	2021									objective, model
KCL [13]	2021									
DT2 [61]	2021									
LTM [46]	2021									
ACE [124]	2021									
ResLT [125]	2021									
SADE [30]	2021									

loss re-weighting by effective number

➤ **Intuition:** Non-overlapping sample number, instead of the vanilla quantity number, playing the role of imbalance



➤ **Effective Number:** The effective number of examples is the expected volume of samples.

$$E_n = (1 - \beta^n) / (1 - \beta)$$

where $\beta = (N - 1) / N$

$$\lim_{\beta \rightarrow 1} E_n = n$$

➤ **Class-Balanced Loss:** Training from imbalanced data by introducing a weighting factor that is **inversely proportional** to the effective number of samples.

The class-balanced loss term can be applied to a wide range of deep networks and loss functions.

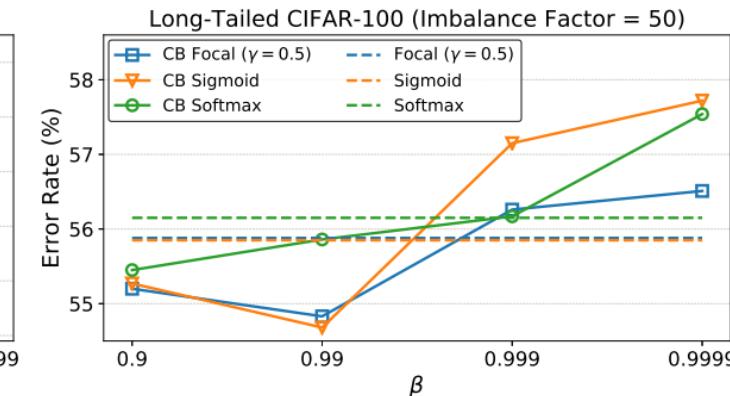
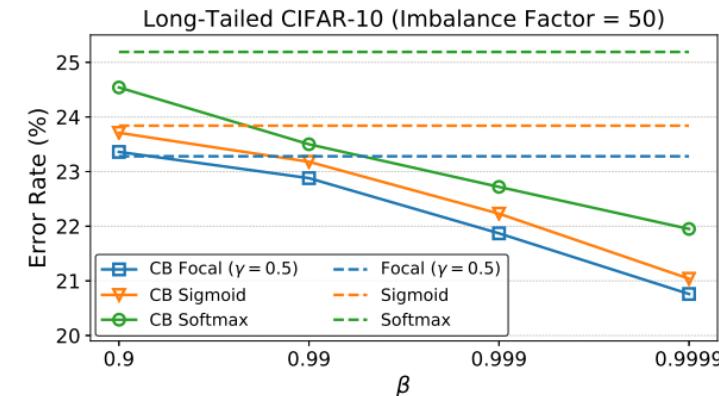
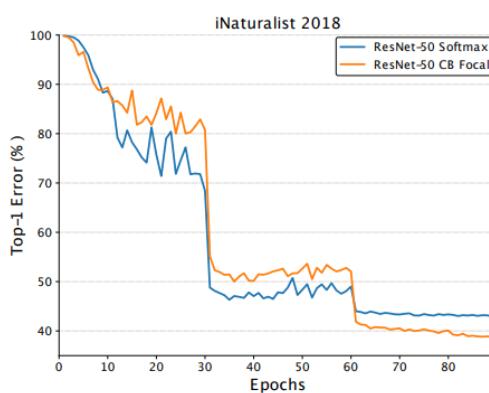
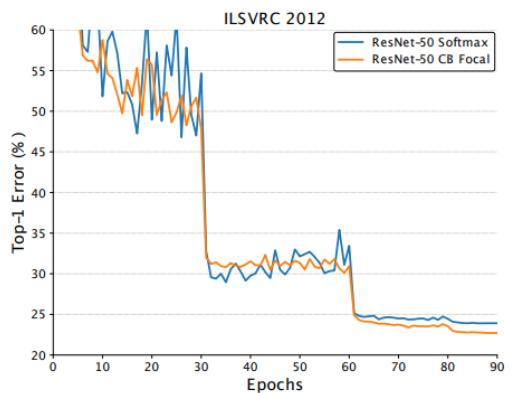
Supervised Long-tailed Learning

➤ **Class-Balanced Loss:** The class-balanced (CB) loss can be written as:

$$\text{CB}(\mathbf{p}, y) = \frac{1}{E_{n_y}} \mathcal{L}(\mathbf{p}, y) = \frac{1 - \beta}{1 - \beta^{n_y}} \mathcal{L}(\mathbf{p}, y)$$

$$\text{CB}_{\text{softmax}}(\mathbf{z}, y) = -\frac{1 - \beta}{1 - \beta^{n_y}} \log \left(\frac{\exp(z_y)}{\sum_{j=1}^C \exp(z_j)} \right)$$

Class-Balanced loss can also be combined with sigmoid cross-entropy loss, focal loss, etc.



Dataset Name	Long-Tailed CIFAR-10						Long-Tailed CIFAR-100					
	200	100	50	20	10	1	200	100	50	20	10	1
Imbalance												
Softmax	34.32	29.64	25.19	17.77	13.61	6.61	65.16	61.68	56.15	48.86	44.29	29.07
Sigmoid	34.51	29.55	23.84	16.40	12.97	6.36	64.39	61.22	55.85	48.57	44.73	28.39
Focal ($\gamma = 0.5$)	36.00	29.77	23.28	17.11	13.19	6.75	65.00	61.31	55.88	48.90	44.30	28.55
Focal ($\gamma = 1.0$)	34.71	29.62	23.29	17.24	13.34	6.60	64.38	61.59	55.68	48.05	44.22	28.85
Focal ($\gamma = 2.0$)	35.12	30.41	23.48	16.77	13.68	6.61	65.25	61.61	56.30	48.98	45.00	28.52
Class-Balanced	31.11	25.43	20.73	15.64	12.51	6.36*	63.77	60.40	54.68	47.41	42.01	28.39*
Loss Type	SM	Focal	Focal	SM	SGM	SGM	Focal	Focal	SGM	Focal	SGM	-
β	0.9999	0.9999	0.9999	0.9999	0.9999	-	0.9	0.9	0.99	0.99	0.99	-
γ	-	1.0	2.0	-	-	-	1.0	1.0	-	0.5	0.5	-

The proposed framework provides a non-parametric means of quantifying data overlap.

Class-wise margin calibration

➤ **Motivation:** Re-weighting and re-sampling often cause over-fitting. The authors propose to regularize the minority classes more strongly than the frequent classes to **improve the generalization error** of minority classes without sacrificing the model's ability to fit the frequent classes.

➤ **Class-distribution-aware margin trade-off:** Enforcing bigger margins can improve generalization, but may also hurt the margins of the frequent classes.

The generalization error is proportional to the following:

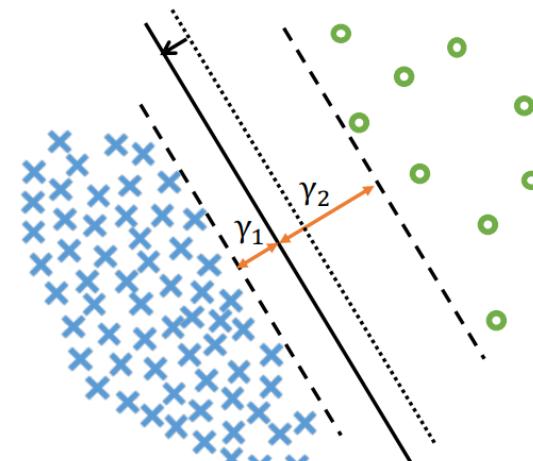
$$\frac{1}{\gamma_1 \sqrt{n_1}} + \frac{1}{\gamma_2 \sqrt{n_2}} \quad \gamma_1 + \gamma_2 = \gamma$$

For multi-class classification tasks, the margin of a sample is defined as:

$$\gamma(x, y) = f(x)_y - \max_{j \neq y} f(x)_j$$

The margin of a class is defined as the minimum value of the spacing between all its samples:

$$\gamma_j = \min_{i \in S_j} \gamma(x_i, y_i) \quad \gamma_j = \frac{C}{n_j^{1/4}}$$



Supervised Long-tailed Learning



➤ **LDAM:** The authors define their loss function as:

$$\mathcal{L}_{\text{LDAM-HG}}((x, y); f) = \max(\max_{j \neq y} \{z_j\} - z_y + \Delta_y, 0)$$

$$\text{where } \Delta_j = \frac{C}{n_j^{1/4}} \text{ for } j \in \{1, \dots, k\}$$

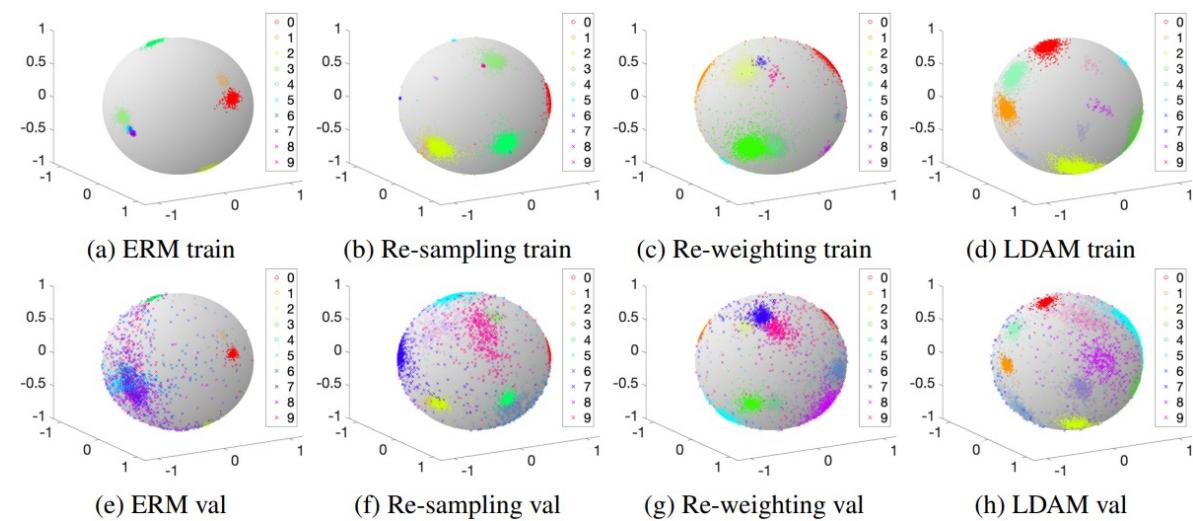
The smooth relaxation of the hinge loss is the following cross-entropy loss with enforced margins:

$$\mathcal{L}_{\text{LDAM}}((x, y); f) = -\log \frac{e^{z_y - \Delta_y}}{e^{z_y - \Delta_y} + \sum_{j \neq y} e^{z_j}}$$

$$\text{where } \Delta_j = \frac{C}{n_j^{1/4}} \text{ for } j \in \{1, \dots, k\}$$

Approach	Error on positive reviews	Error on negative reviews	Mean Error
ERM	2.86	70.78	36.82
RS	7.12	45.88	26.50
RW	5.20	42.12	23.66
LDAM-DRW	4.91	30.77	17.84

Dataset	Imbalanced CIFAR-10				Imbalanced CIFAR-100			
	long-tailed		step		long-tailed		step	
Imbalance Type	100	10	100	10	100	10	100	10
ERM	29.64	13.61	36.70	17.50	61.68	44.30	61.45	45.37
Focal [Lin et al., 2017]	29.62	13.34	36.09	16.36	61.59	44.22	61.43	46.54
LDAM	26.65	13.04	33.42	15.00	60.40	43.09	60.42	43.73
CB RS	29.45	13.21	38.14	15.41	66.56	44.94	66.23	46.92
CB RW [Cui et al., 2019]	27.63	13.46	38.06	16.20	66.01	42.88	78.69	47.52
CB Focal [Cui et al., 2019]	25.43	12.90	39.73	16.54	63.98	42.01	80.24	49.98
HG-DRS	27.16	14.03	29.93	14.85	-	-	-	-
LDAM-HG-DRS	24.42	12.72	24.53	12.82	-	-	-	-
M-DRW	24.94	13.57	27.67	13.17	59.49	43.78	58.91	44.72
LDAM-DRW	22.97	11.84	23.08	12.19	57.96	41.29	54.64	40.54





Supervised Long-tailed Learning



Class-wise logit adjustment

- **Motivation:** Design a consistent loss function that allows for a relatively **elastic margin** in the logit for head and tail.
- **Balanced error:** Under class imbalance, to measure balanced error:

$$\text{BER}(f) \doteq \frac{1}{L} \sum_{y \in [L]} \mathbb{P}_{x|y} \left(y \notin \operatorname{argmax}_{y' \in \mathcal{Y}} f_{y'}(x) \right)$$

Under Bayes-optimal prediction, if $\mathbb{P}^{\text{bal}}(y | x) \propto \mathbb{P}(y | x) / \mathbb{P}(y)$

Then

$$\operatorname{argmax}_{y \in [L]} \mathbb{P}^{\text{bal}}(y | x) = \operatorname{argmax}_{y \in [L]} \exp(s_y^*(x)) / \mathbb{P}(y) = \operatorname{argmax}_{y \in [L]} s_y^*(x) - \ln \mathbb{P}(y)$$





Supervised Long-tailed Learning



➤ The logit adjusted softmax cross-entropy

$$\ell(y, f(x)) = -\log \frac{e^{f_y(x) + \tau \cdot \log \pi_y}}{\sum_{y' \in [L]} e^{f_{y'}(x) + \tau \cdot \log \pi_{y'}}} = \log \left[1 + \sum_{y' \neq y} \left(\frac{\pi_{y'}}{\pi_y} \right)^\tau \cdot e^{(f_{y'}(x) - f_y(x))} \right]$$

$$w_1^\top \Phi(x)/\pi_1 < w_2^\top \Phi(x)/\pi_2 \not\iff \exp(w_1^\top \Phi(x))/\pi_1 < \exp(w_2^\top \Phi(x))/\pi_2.$$

➤ Post-hoc logit adjustment

$$\operatorname{argmax}_{y \in [L]} \exp(w_y^\top \Phi(x)) / \pi_y^\tau = \operatorname{argmax}_{y \in [L]} f_y(x) - \tau \cdot \log \pi_y$$





Supervised Long-tailed Learning



➤ An remarkable point on the statistical consistency of long-tailed multi-class classification

$$\ell(y, f(x)) = \alpha_y \cdot \log \left[1 + \sum_{y' \neq y} e^{\Delta_{yy'}} \cdot e^{(f_{y'}(x) - f_y(x))} \right]$$

↑

Theorem 1. For any $\delta \in \mathbb{R}_+^L$, the pairwise loss in (11) is Fisher consistent with weights and margins

$$\alpha_y = \delta_y / \mathbb{P}(y) \quad \Delta_{yy'} = \log (\delta_{y'} / \delta_y).$$

Letting $\delta_y = \pi_y$, we immediately deduce that the logit-adjusted loss of (10) is consistent, provided our π_y is a consistent estimate of $\mathbb{P}(y)$. Similarly, $\delta_y = 1$ recovers the classic result that the balanced loss is consistent. While Theorem 1 only provides a sufficient condition in multi-class setting, one can provide a necessary and sufficient condition that rules out other choices of Δ in the binary case.





Supervised Long-tailed Learning



Dynamic adjustment based on a fine-grained generalization bound

Proposition 3 (Data-Dependent Bound for the VS Loss). *Given the function set \mathcal{F} and the VS loss L_{VS} , for any $\delta \in (0, 1)$, with probability at least $1 - \delta$ over the training set \mathcal{S} , the following generalization bound holds for all $f \in \mathcal{F}$:*

$$\mathcal{R}_{bal}^L(f) \lesssim \Phi(L_{VS}, \delta) + \frac{\hat{\mathfrak{C}}_{\mathcal{S}}(\mathcal{F})}{C\pi_C} \sum_{y=1}^C \alpha_y \tilde{\beta}_y \sqrt{\pi_y} [1 - \text{softmax}(\beta_y B_y(f) + \Delta_y)].$$

$$L_{VS}(f(\mathbf{x}), y) = -\alpha_y \log \left(\frac{e^{\beta_y f(\mathbf{x})_y + \Delta_y}}{\sum_{y'} e^{\beta_{y'} f(\mathbf{x})_{y'} + \Delta_{y'}}} \right).$$

Algorithm 1: Principled Learning Algorithm induced by the Theoretical Insights

Require: Training set $\mathcal{S} = \{(x_i, y_i)\}_{i=1}^N$ and a model f parameterized by Θ .

- 1: Initialize the model parameters Θ randomly.
 - 2: **for** $t = 1, 2, \dots, T$ **do**
 - 3: $\mathcal{B} \leftarrow \text{SampleMiniBatch}(\mathcal{S}, m)$ ▷ A mini-batch of m samples
 - 4: **if** $t < T_0$ **then**
 - 5: Set $\alpha = 1, \beta_y, \Delta_y$ ▷ Adjust logits during the initial phase
 - 6: **else**
 - 7: Set $\alpha_y \propto \pi_y^{-\nu}, \beta_y = 1, \Delta_y, \nu > 0$ ▷ TLA and ADRW
 - 8: **end if**
 - 9: $L(f, \mathcal{B}) \leftarrow \frac{1}{m} \sum_{(\mathbf{x}, y) \in \mathcal{B}} L_{VS}(f(\mathbf{x}), y)$ ▷ Calculate the loss
 - 10: $\Theta \leftarrow \Theta - \eta \nabla_{\Theta} L(f, \mathcal{B})$ ▷ One SGD step
 - 11: Optional: anneal the learning rate η . ▷ Required when $t = T_0$
 - 12: **end for**
-



Supervised Long-tailed Learning

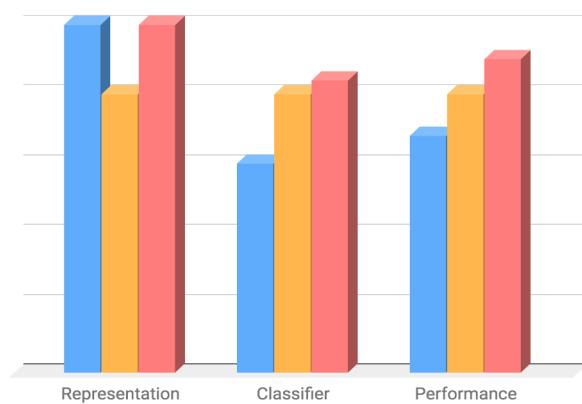


Both representation and classifier matter

➤ **Motivation:** Representation and classifier are in the different learning pace, and can be treated differently during training.

➤ **The problem behind long-tail:**

Classification performance =
Representation Quality + Classifier Quality



➤ **Ways of learning classifiers:**

Classifier Re-training (cRT)

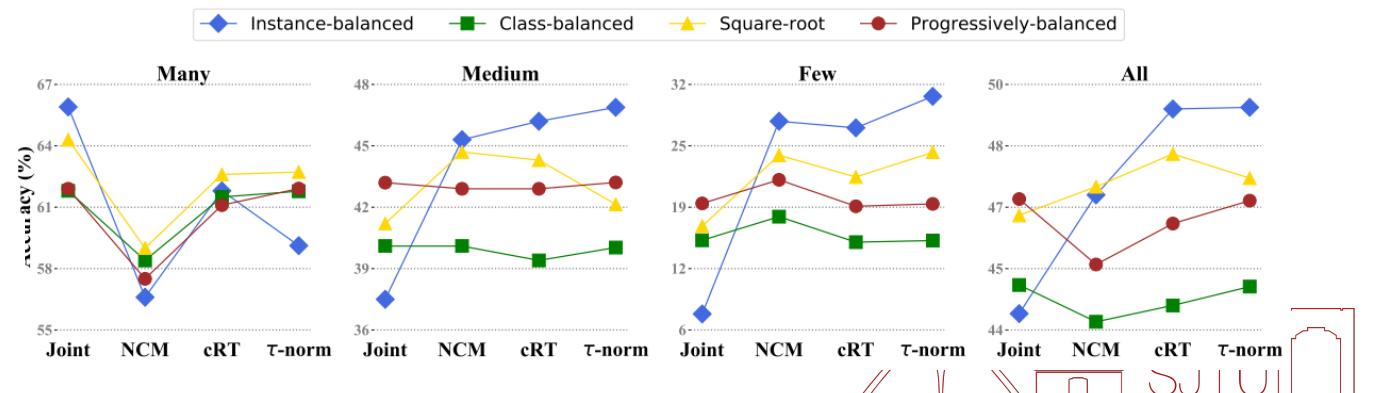
Nearest Class Mean classifier (NCM)

τ -normalized classifier (τ -normalized)

$$\widetilde{w}_i = \frac{w_i}{\|w_i\|^\tau}$$

Learnable weight scaling (LWS)

$$\widetilde{w}_i = f_i * w_i, \text{ where } f_i = \frac{1}{\|w_i\|^\tau}$$





Supervised Long-tailed Learning



The special constant classifier geometry

➤ **Motivation:** In the training on an imbalanced dataset, the classifier vectors of minor classes will be merged, termed as minority collapse, which breaks up the ETF structure and deteriorates the performance on test data.

➤ **Simplex Equiangular Tight Frame:**

$$\mathbf{M} = \sqrt{\frac{K}{K-1}} \mathbf{U} \left(\mathbf{I}_K - \frac{1}{K} \mathbf{1}_K \mathbf{1}_K^T \right)$$
$$\mathbf{m}_i^T \mathbf{m}_j = \frac{K}{K-1} \delta_{i,j} - \frac{1}{K-1}, \forall i, j \in [1, K]$$

➤ **ETF Classifier:**

$$\min_{\mathbf{H}} \quad \frac{1}{N} \sum_{k=1}^K \sum_{i=1}^{n_k} \mathcal{L}_{CE}(\mathbf{h}_{k,i}, \mathbf{W}^*)$$
$$\mathbf{w}_k^{*T} \mathbf{w}_{k'}^* = E_W \left(\frac{K}{K-1} \delta_{k,k'} - \frac{1}{K-1} \right), \forall k, k' \in [1, K]$$



Supervised Long-tailed Learning

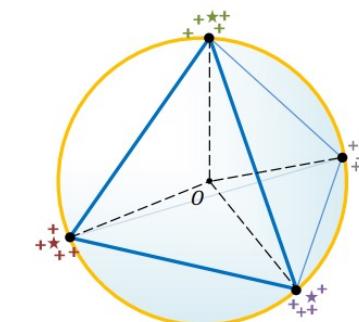
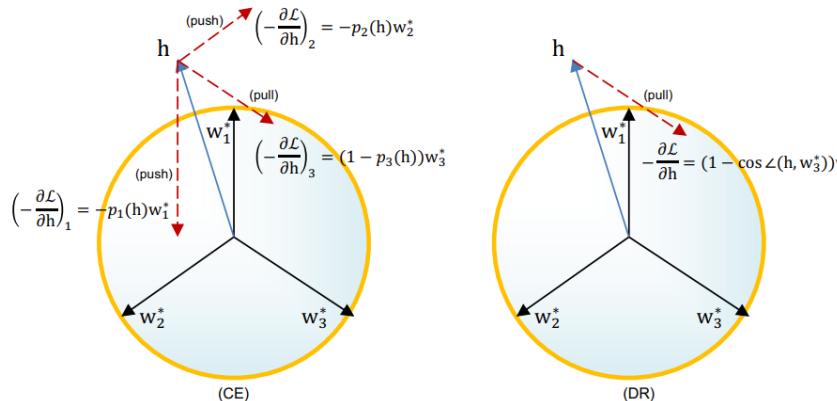


➤ Dot-Regression Loss:

Squared loss function:

$$\mathcal{L}_{DR}(\mathbf{h}, \mathbf{W}^*) = \frac{1}{2\sqrt{E_W E_H}} \left(\mathbf{w}_c^{*T} \mathbf{h} - \sqrt{E_W E_H} \right)^2$$

$$\frac{\partial \mathcal{L}_{DR}}{\partial \mathbf{h}} = - (1 - \cos \angle(\mathbf{h}, \mathbf{w}_c^*)) \mathbf{w}_c^*$$



Epoch	Methods	Acc. (%)
90	Learnable Classifier + CE	34.6
	ETF Classifier + DR	41.8
120	Learnable Classifier + CE	41.9
	ETF Classifier + DR	43.2
150	Learnable Classifier + CE	42.5
	ETF Classifier + DR	43.8
180	Learnable Classifier + CE	44.3
	ETF Classifier + DR	44.7

Backbone	Methods	Acc. (%)
ResNet-34	Learnable Classifier + CE	82.2
	ETF Classifier + DR	83.0
ResNet-50	Learnable Classifier + CE	85.5
	ETF Classifier + DR	86.1
ResNet-101	Learnable Classifier + CE	86.2
	ETF Classifier + DR	87.0

Methods	CIFAR-10 [17]			CIFAR-100 [17]			SVHN [25]			STL-10 [3]		
	0.005	0.01	0.02	0.005	0.01	0.02	0.005	0.01	0.02	0.005	0.01	0.02
<i>ResNet</i>												
Learnable Classifier + CE	67.3	72.8	78.6	38.7	43.0	48.1	40.5	40.9	49.3	33.1	37.9	38.8
ETF Classifier + DR	71.9	76.5	81.0	40.9	45.3	50.4	42.8	45.7	49.8	33.5	37.2	37.9
Improvements	+4.6	+3.7	+2.4	+2.2	+2.3	+2.3	+2.3	+4.8	+0.5	+0.4	-0.7	-0.9
<i>DenseNet</i>												
Learnable Classifier + CE	71.1	77.7	84.1	40.3	43.8	49.8	39.7	40.5	46.4	38.5	41.2	44.9
ETF Classifier + DR	72.9	78.5	83.4	40.1	44.0	49.7	40.5	44.8	48.4	39.5	42.9	46.3
Improvements	+1.8	+0.8	-0.7	-0.2	+0.2	-0.1	+0.8	+4.3	+2.0	+1.0	+1.7	+1.4

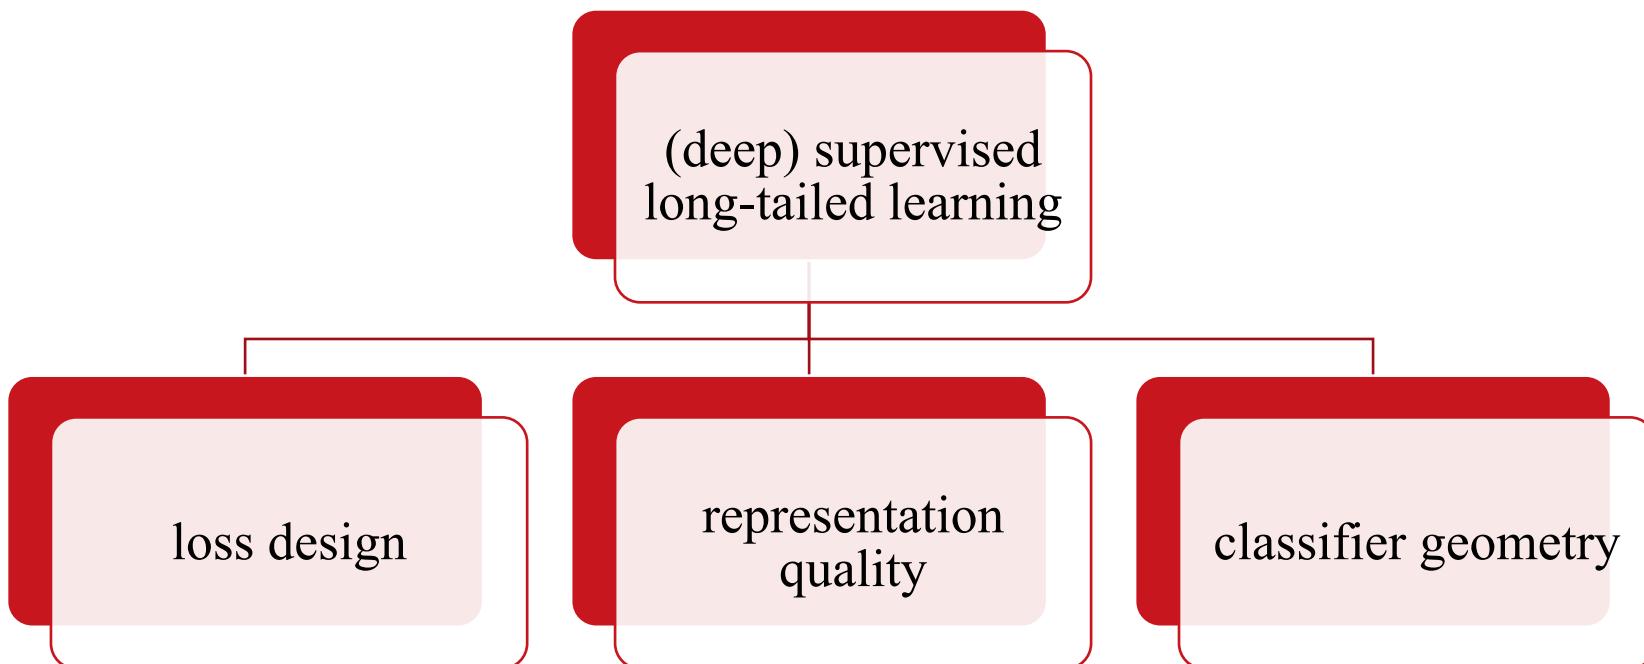
[1] Yang et al. "Inducing Neural Collapse in Imbalanced Learning: Do We Really Need a Learnable Classifier at the End of Deep Neural Network?" NeurIPS 2022.



Supervised Long-tailed Learning



Summary: supervised long-tailed learning in the context of deep learning, should take more factors (loss, representation etc.) into account to improve the performance of imbalanced learning.



Weak supervision

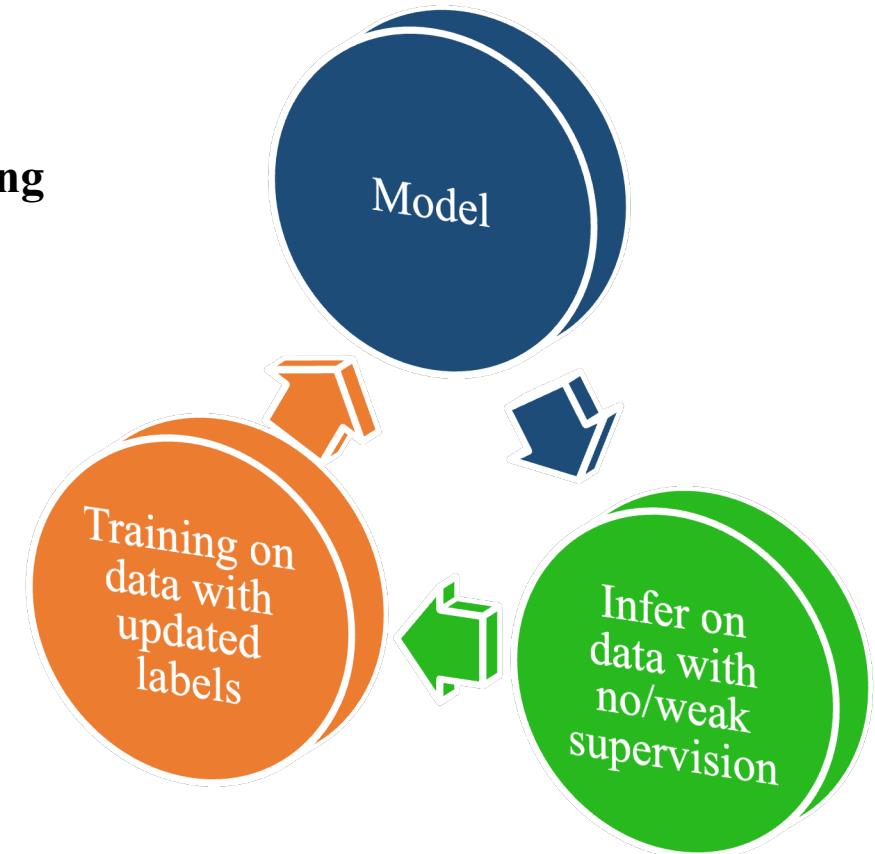
- Incomplete supervision: **semi-supervised learning**
- Inexact supervision: **partial label learning, multi-instance learning**
- Inaccurate supervision: **label-noise learning**

Target: *low labeling cost but high accuracy.*

Pseudo Labeling

- Generate more precise annotations for weakly supervised data

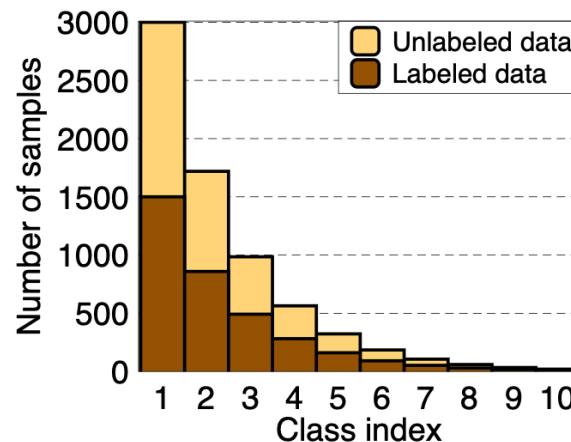
Transfer weakly-supervised learning → supervised learning



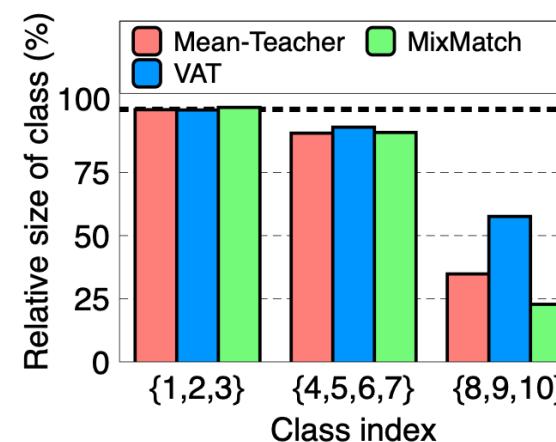
Pseudo Labeling

When long-tailed distribution meets weak supervision

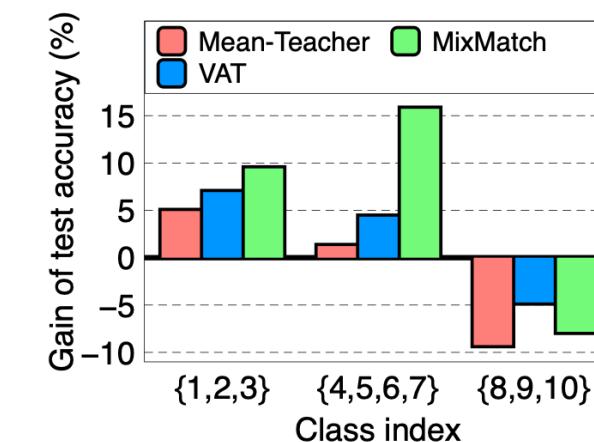
- Positive feedback property of pseudo-labeling: head get better and more



(a) Imbalanced class distribution



(b) Bias of pseudo-labels



[1]

(c) Accuracy gain from SSL

- Distributional corrections need to go deeper into training and pseudo-labeling





Long-tailed Semi-Supervised Learning



Long-tailed Semi-Supervised Learning

- Motivation: Label propagation on **unlabeled data biased toward the majority classes**.
- Main idea: **Refine** the original, biased pseudo-labels so that their distribution can **match the true class distribution of unlabeled data**, while constrain the refined pseudo-labels to be close from the original ones.
- Formulation: An optimization problem for constructing refined pseudo-labels: **minimizing the distortion from the original pseudo-labels, while matching the true class distribution**.

$$\text{minimize} \quad \sum_{m=1}^M w_m D_{KL}(\hat{y}_m \parallel \hat{y}_m^{\text{unlabeled}})$$

$$\text{subject to} \quad \sum_{m=1}^M \hat{y}_m(k) = M_k, \quad \forall k, \quad \sum_{k=1}^K \hat{y}_m(k) = 1, \quad \forall m, \quad \hat{y}_m(k) \in [0, 1], \quad \forall m, k$$

Algorithm 1 DualCoordinateAscent: Coordinate ascent algorithm for dual of (1)

Require: $\{\hat{y}_m^{\text{unlabeled}}\}_{m=1}^M, \{w_m\}_{m=1}^M, \{M_k\}_{k=1}^K, T$

Ensure: The unique solution of (1)

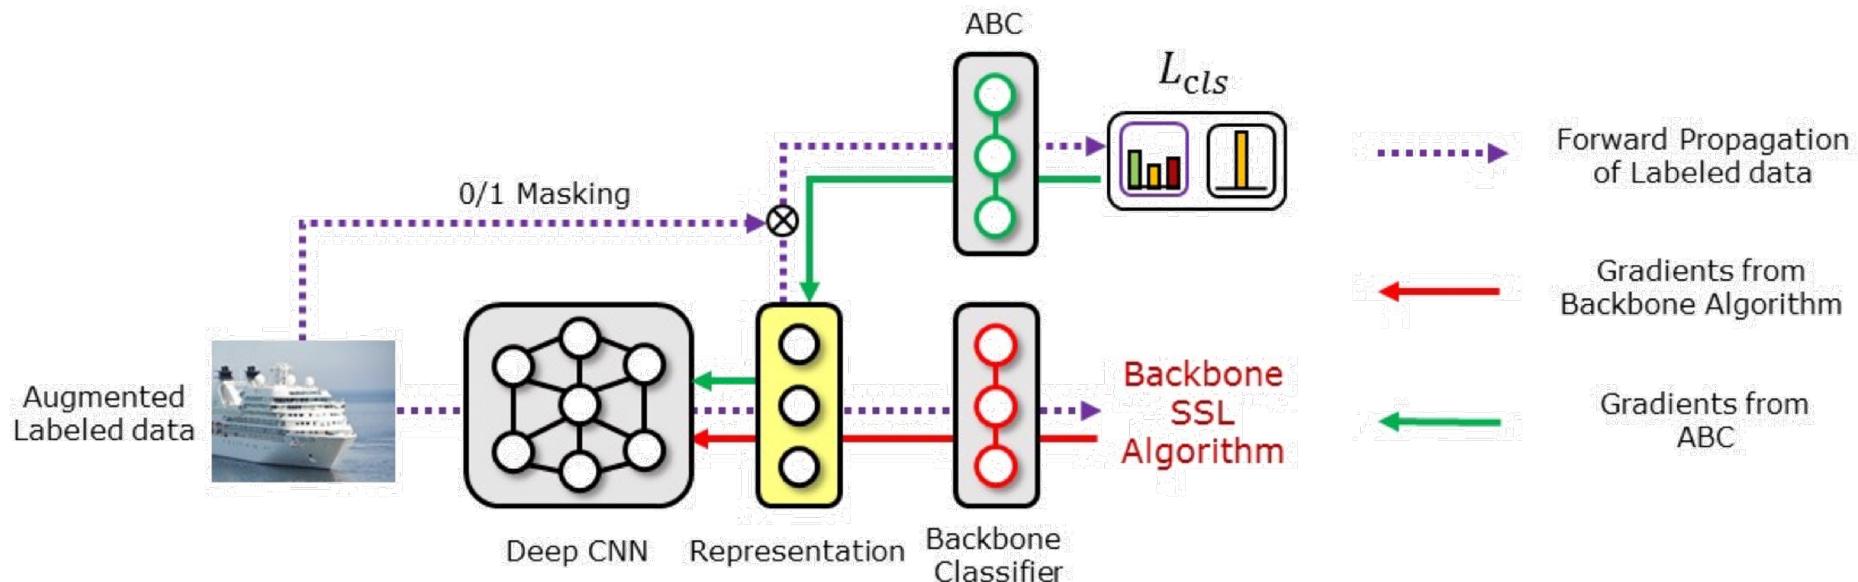
```

1:  $\hat{y}_m^0 \leftarrow y_m^{\text{unlabeled}}$ ,  $\alpha_m^0 \leftarrow 1$ ,  $\beta_k^0 \leftarrow 1$ ,  $\forall m, k$ 
2: for  $t = 1$  to  $T$  do
3:   if  $t$  is odd or  $t = T$  then
4:      $\beta_k^t \leftarrow \beta_k^{t-1}$ ,  $\alpha_m^t \leftarrow \left( \sum_{k=1}^K \hat{y}_m^0(k) (\beta_k^{t-1})^{\frac{1}{w_m}} \right)^{-1}$ ,  $\forall m, k$ ,
5:   else
6:      $\alpha_m^t \leftarrow \alpha_m^{t-1}$ ,  $\beta_k^t \leftarrow \text{Solve}_{Z \geq 0} \left( \sum_{m=1}^M \hat{y}_m^0(k) \alpha_m^{t-1} Z^{\frac{1}{w_m}} - M_k \right)$ ,  $\forall m, k$ 
7:   end if
8: end for
9:  $\hat{y}_m^{\text{out}}(k) \leftarrow \hat{y}_m^0(k) \alpha_m^T (\beta_k^T)^{\frac{1}{w_m}}$ ,  $\forall m, k$ 

```

ABC: Auxiliary Balanced Classifier [1]

- Motivation: high-quality representations can be learned even the classifier is biased [2].
- Train an Auxiliary Balanced Classifier (ABC) by resampling a balanced subset while using the high-quality representations learned from all data.



- [1] Lee, H., Shin, S., & Kim, H.. Abc: Auxiliary balanced classifier for class-imbalanced semi-supervised learning. *NeurIPS*. 2021.
[2] Kang, B., Xie, S., Rohrbach, M., et al. Decoupling Representation and Classifier for Long-Tailed Recognition. *ICLR*. 2020.

Long-tailed Partial Label Learning

- Dataset $\mathcal{X} = \{(x_i, y_i, S_i) : i \in (1, 2, 3, \dots, N)\}$, where any x_i is associated with a candidate label set S_i and its ground truth $y_i \in S_i$ is invisible.



Parker: Morning, Flash.
Thompson: Good morning, Parker.



The History of LeBron James and Stephen Curry's Rivalrous Friendship



Annotator 1: Korat
Annotator 2: Russian Blue

- The sample number of L classes in descending order

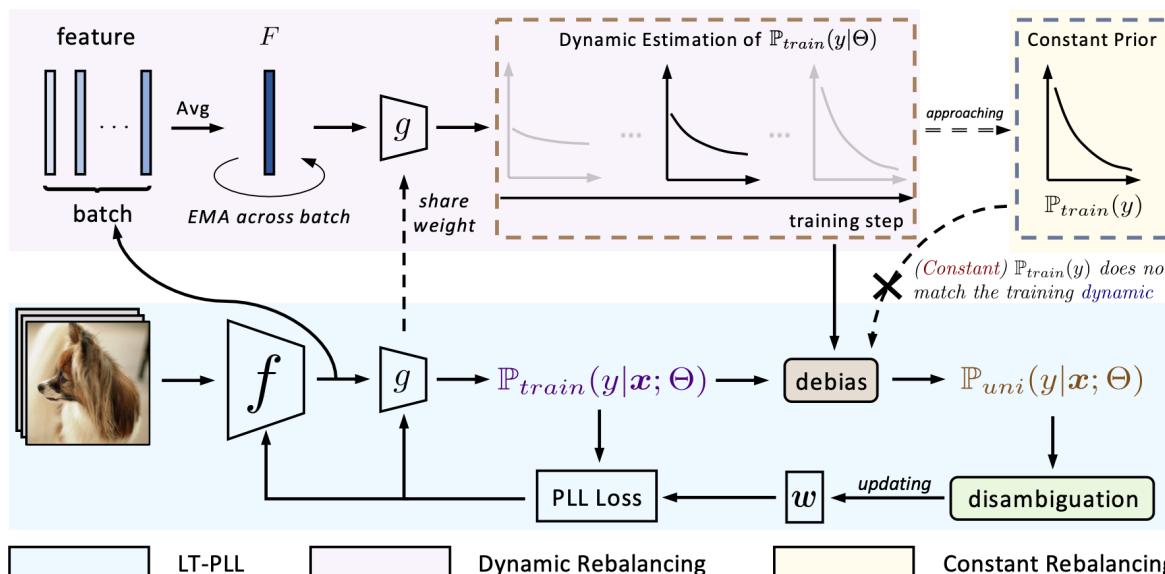
$$\forall N_1 \geq N_2 \geq \dots \geq N_L, \text{Imbalance ratio } \gamma = \frac{N_1}{N_L} \gg 1$$

- Main challenges:

- Tail samples cannot be correctly recognized even in training
- No available class prior

RECORDS: Rebalancing for dynamic bias

- Observation: even after applying an oracle class distribution prior in the training, existing long-tailed techniques underperform in LT-PLL and even fail in some cases.
 - Existing long-tailed techniques: leverage a constant class distribution prior to rebalance the training and does not consider the dynamic of label disambiguation.
- A dynamic rebalancing mechanism friendly to the training dynamic



Constant Rebalancing

$$z_{uni}^y(\mathbf{x}) = z^y(\mathbf{x}) - \log \mathbb{P}_{train}(y|\Theta) = z^y(\mathbf{x}) - \log \text{softmax}(g^y(F; \mathbf{W})).$$

Match the training dynamic

Dynamic Rebalancing

$$z_{uni}^y(\mathbf{x}) = z^y(\mathbf{x}) - \log \mathbb{P}_{train}(y|\Theta) = z^y(\mathbf{x}) - \log \text{softmax}(g^y(F; \mathbf{W})).$$

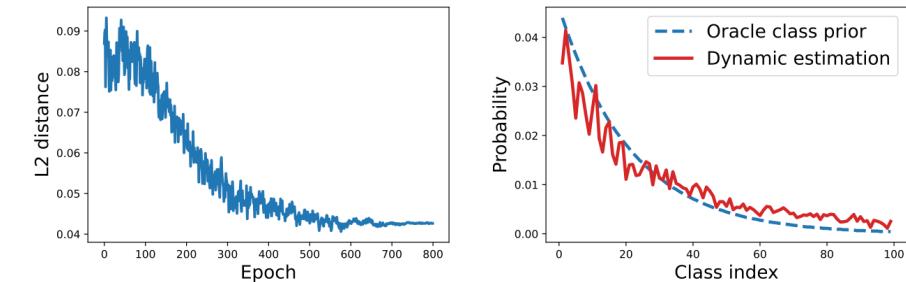
RECORDS: Rebalancing for dynamic bias

- Theory: alongside the label disambiguation, the dynamic estimation can progressively approach to the oracle class distribution.

Proposition 1. Let $\eta = \sup_{(\mathbf{x}, y) \in \mathcal{X} \times \mathcal{Y}, y_j \in \mathcal{Y}, y_j \neq y} \mathbb{P}_{S|(\mathbf{x}, y)}(y_j \in S)$ denote the ambiguity degree, d_H be the Natarajan dimension of the hypothesis space H , $\tilde{h} = h_{\tilde{\Theta}}$ be the optimal classifier on the basis of the label disambiguation, where $\tilde{\Theta} = \arg \min_{\Theta} R^{\mathcal{D}_{train}}(\Theta)$. If the small ambiguity degree condition (Cour et al., 2011a; Liu & Dietterich, 2014)) satisfies, namely, $\eta \in [0, 1)$, then for $\forall \delta > 0$, the L_2 distance between $\mathbb{P}_{train}(y)$ and $\mathbb{P}_{train}(y|\tilde{\Theta})$ given \tilde{h} is bounded as

$$L_2(\tilde{h}) < \frac{4}{(\ln 2 - \ln(1 + \eta))N} (d_H(\ln 2N + 2 \ln C) - \ln \delta + \ln 2)$$

with probability at least $1 - \delta$, where N is the sample number and C is the category number.



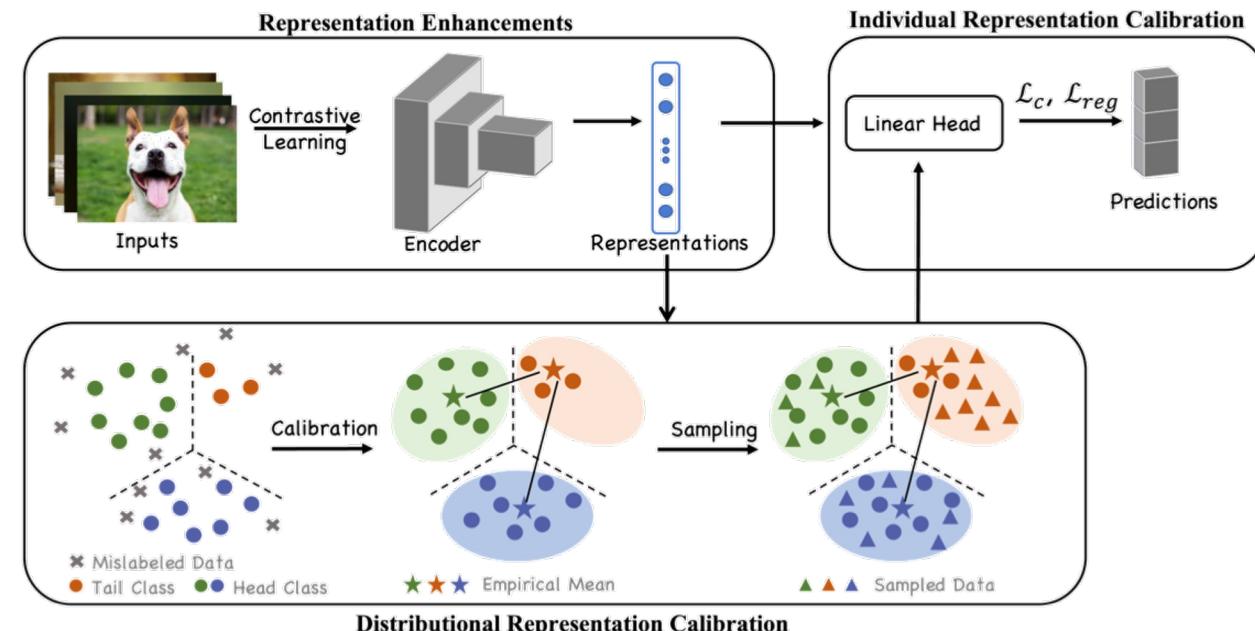
(a) L_2 distance during training

(b) Estimated class distribution

	CIFAR-10-LT						CIFAR-100-LT						PASCAL VOC
Imbalance ratio ρ	50			100			50			100			
Ambiguity q	0.3	0.5	0.7	0.3	0.5	0.7	0.03	0.05	0.07	0.03	0.05	0.07	
LW	71.31	39.69	25.23	66.08	41.2	26.59	37.34	31.59	30.6	33.78	30.69	27.91	20.73
CAVL	57.77	41.68	21.32	54.28	39.47	20.67	23.01	21.29	14.23	31.65	28.77	13.65	18.33
PRODEN	74.22	56.25	45.11	65.72	49.23	42.26	43.02	40.15	39.1	38.66	36.23	35.52	24.47
CORR	77.53	58.35	45.66	69.12	50.96	42.87	45.95	42.83	41.72	41.79	38.47	36.96	26.72
SoLar	83.88	76.55	54.61	75.38	70.63	53.15	47.93	46.85	45.1	42.51	41.71	39.15	56.49
LW+RECORDS	77.22	59.88	43.98	72.2	59.87	43.63	38.16	33.6	31.53	35.33	31.55	28.8	55.19
CAVL+RECORDS	68.93	63.59	44.61	66.65	59.62	41.88	46.13	41.93	34.01	41.33	35.79	31.29	55.03
PRODEN+RECORDS	81.23	78.82	68.03	73.92	66.21	55.73	48.12	46.01	44.19	42.98	41.03	40.25	54.78
CORR+RECORDS	84.25	82.5	71.24	79.79	74.07	62.25	52.08	50.58	47.91	46.57	45.22	44.73	58.45

RCAL: Representation Calibration

- Motivation: The representations of **unsupervised contrastive learning** are not influenced by corrupted labels and thus **naturally robust**.
- Main spirit: **Decoupling the imbalance from the noise labels** with the help of self-supervised representation learning methods.
- Based upon the achieved representations, the calibration can be performed. **Distributional representation calibration** estimate the robust feature distribution considering data imbalance and perform balanced sampling. **Individual representation calibration** constrain the distance between supervised training representations and pre-training representations.



RCAL: Representation Calibration

Algorithm 1 Algorithm of the proposed method RCAL

Require: the training dataset $\tilde{\mathcal{S}} = \{(\mathbf{x}_i, \tilde{y}_i)\}_{i=1}^n$, regularization strength β , scalar temperature τ , confidence weight γ , the pre-training epochs T_p , max epochs T_m .

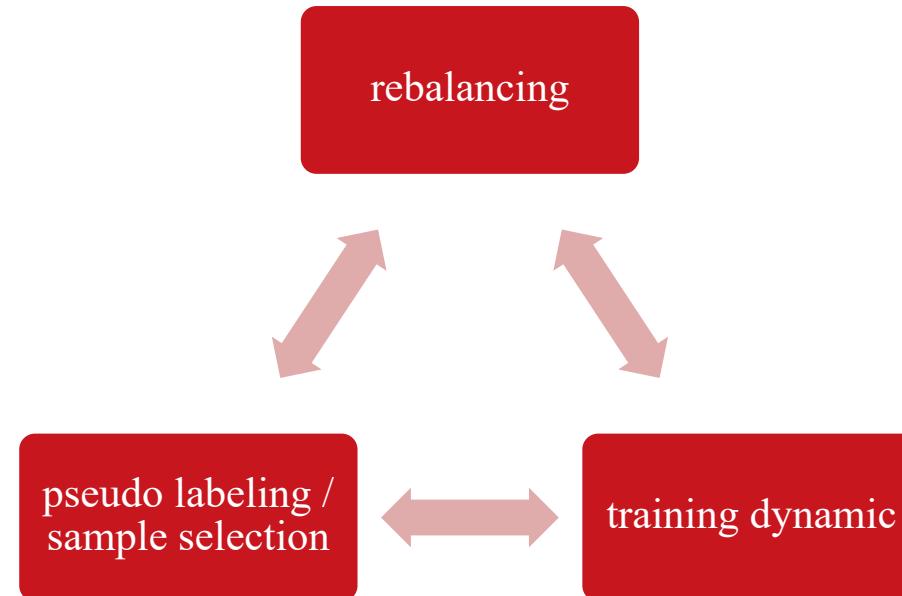
- 1: **for** $t = 1, \dots, T_p$ **do**
- 2: **Pre-train** the encoder network f with MoCo [20].
- 3: **end for**
- 4: **Extract** deep representations of instances with $\mathbf{z} = f(\mathbf{x})$.
- 5: **for** $c = 1, \dots, K$ **do**
- 6: **Perform** the LOF algorithm for the c -th class and obtain preserved examples $\tilde{\mathcal{S}}'_c$.
- 7: **Build** the multivariate Gaussian distribution $\mathcal{N}(f(\mathbf{x}) | \hat{\mu}_c, \hat{\Sigma}_c)$ for c -th class using $\tilde{\mathcal{S}}'_c$.
- 8: **end for**
- 9: **Calibrate** the multivariate Gaussian distributions of tail classes with the statistics of head classes.
- 10: **Sample** data points from achieved multivariate Gaussian distributions of all classes.
- 11: **for** $t = T_p + 1, \dots, T_m$ **do**
- 12: **Add** distance constraints between learned representations and representations brought by contrastive learning.
- 13: **Adopt** the mixup technology to original examples.
- 14: **Train** the encoder f and the linear head h simultaneously on the training dataset and sample data points with the training loss in Eq. (2).
- 15: **end for**
- 16: **return** The robust classifier $h(f(\mathbf{x}))$ for testing.

Dataset	Imbalance Ratio	10					100				
		Noise Rate	0.1	0.2	0.3	0.4	0.5	0.1	0.2	0.3	0.4
CIFAR-10	ERM	80.41	75.61	71.94	70.13	63.25	64.41	62.17	52.94	48.11	38.71
	ERM-DRW	81.72	77.61	71.94	70.13	63.25	66.74	62.17	52.94	48.11	38.71
	LDAM	84.59	82.37	77.48	71.41	60.30	71.46	66.26	58.34	46.64	36.66
	LDAM-DRW	85.94	83.73	80.20	74.87	67.93	76.58	72.28	66.68	57.51	43.23
	CRT	80.22	76.15	74.17	70.05	64.15	61.54	59.52	54.05	50.12	36.73
	NCM	82.33	74.73	74.76	68.43	64.82	68.09	66.25	60.91	55.47	42.61
	MiSLAS	87.58	85.21	83.39	76.16	72.46	75.62	71.48	67.90	62.04	54.54
	Co-teaching	80.30	78.54	68.71	57.10	46.77	55.58	50.29	38.01	30.75	22.85
	CDR	81.68	78.09	73.86	68.12	62.24	60.47	55.34	46.32	42.51	32.44
	Sel-CL+	86.47	85.11	84.41	80.35	77.27	72.31	71.02	65.70	61.37	56.21
CIFAR-100	HAR-DRW	84.09	82.43	80.41	77.43	67.39	70.81	67.88	48.59	54.23	42.80
	RoLT	85.68	85.43	83.50	80.92	78.96	73.02	71.20	66.53	57.86	48.98
	RoLT-DRW	86.24	85.49	84.11	81.99	80.05	76.22	74.92	71.08	63.61	55.06
	RCAL (Ours)	88.09	86.46	84.58	83.43	80.80	78.60	75.81	72.76	69.78	65.05

Dataset	Imbalance Ratio	CIFAR-10				CIFAR-100			
		10	100	10	100	10	100	10	100
Noise Rate		0.2	0.4	0.2	0.4	0.2	0.4	0.2	0.4
RCAL		86.46	83.43	75.81	69.78	54.85	48.91	39.85	33.36
RCAL w/o Mixup		84.08	79.27	72.47	64.83	51.22	45.53	36.78	30.85
RCAL w/o Mixup, REG		83.23	78.12	67.49	58.27	48.74	42.15	34.31	27.14
RCAL w/o Mixup, REG, DC		80.40	74.37	64.02	54.61	47.01	40.85	32.27	25.42
RCAL w/o Mixup, REG, DC, CL		75.61	70.13	62.17	48.11	43.27	32.94	26.21	17.91

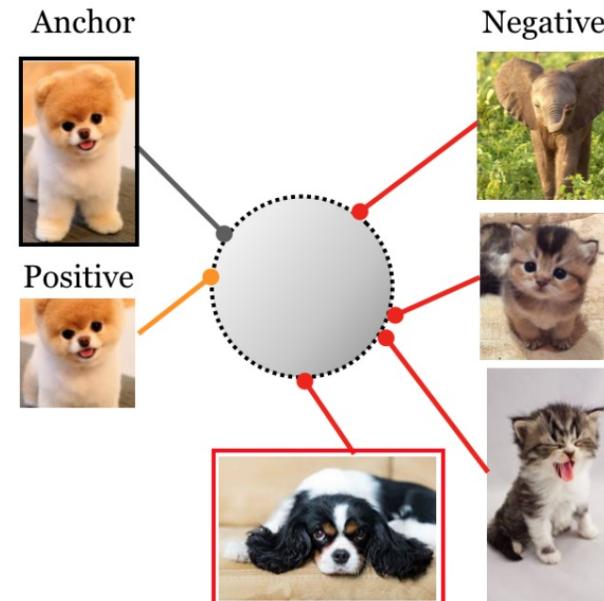
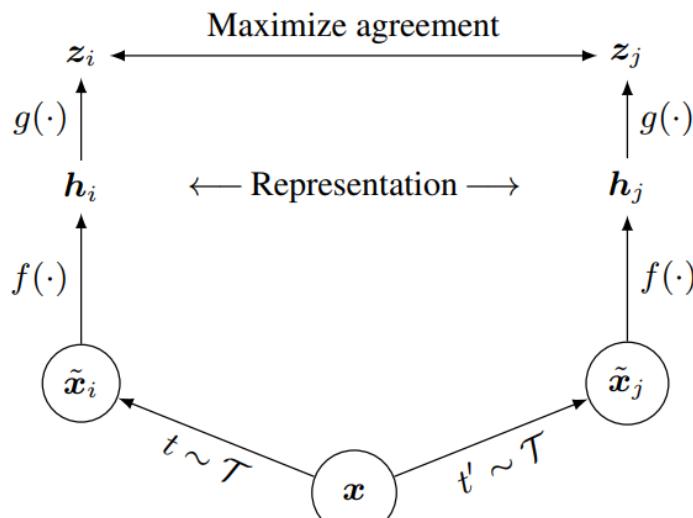
Summary

- Weak Supervision can **exacerbate the long-tail effect**, impairing the performance of minority classes.
- It is critical to construct a friendly mechanism to decouple *pseudo labeling / sample selection* and *rebalancing*.



Self-supervised learning

- The idea of contrastive learning: push augmented views of the same image closer, and embeddings to random pairs of images far apart

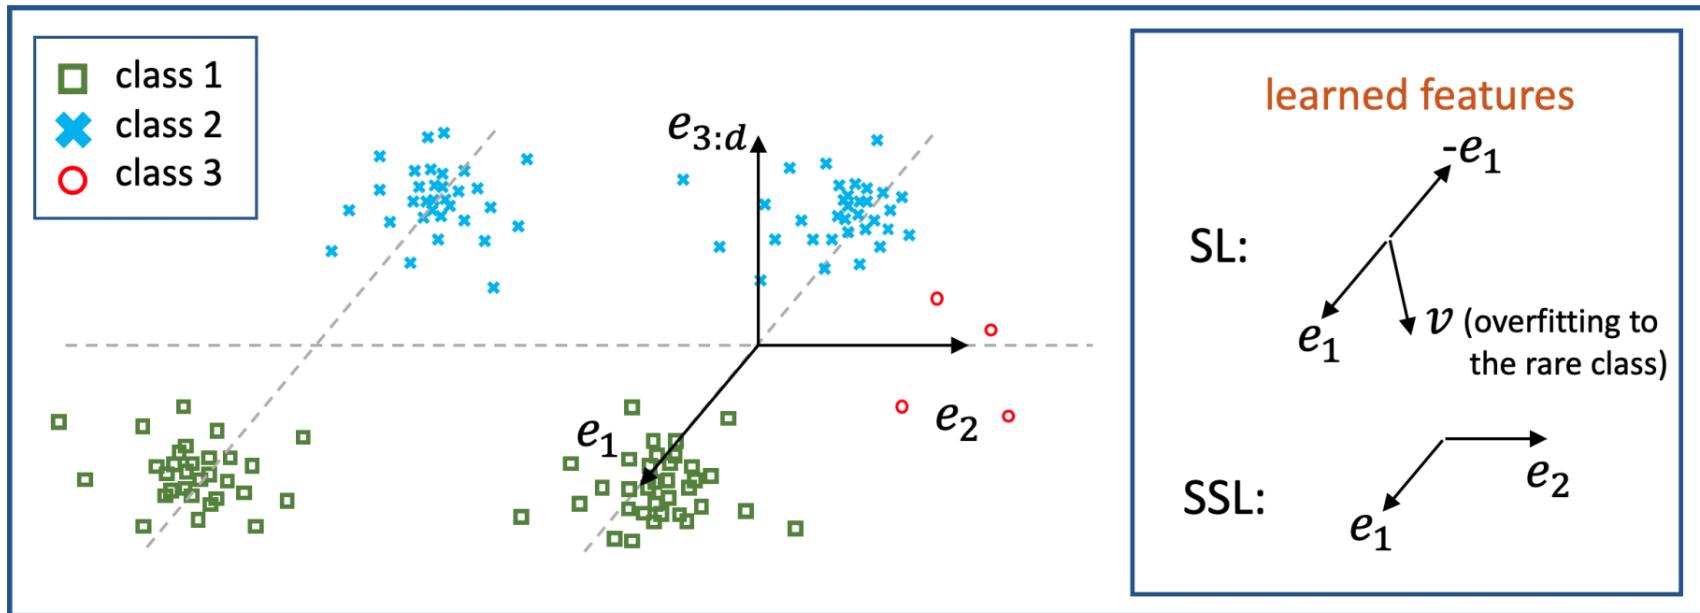


- [1] Chen et al. "A simple framework for contrastive learning of visual representations." ICML 2020.
- [2] Khosla et al. "Supervised contrastive learning." NeurIPS 2020.

Self-Supervised Long-tailed Learning



Is self-supervised learning more robust to data imbalance?



- Supervised learning (SL) only **extracts features that are useful for predicting labels** (e_1)
- Self-supervised learning (SSL) learns **task-irrelevant features regardless of the labels**, which enables richer and more robust representation (e_1, e_2)

Self-Supervised Long-tailed Learning

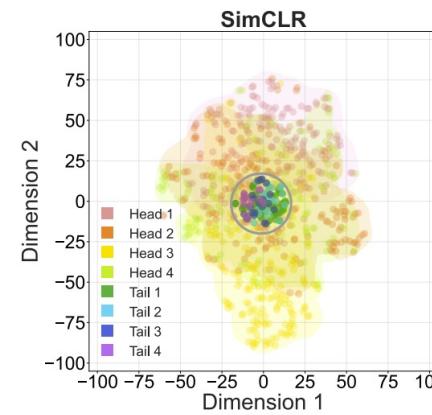
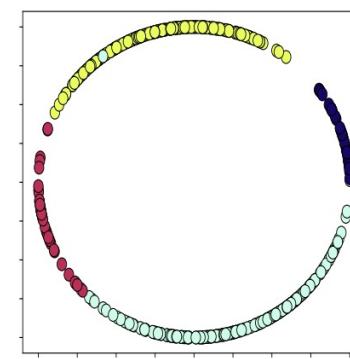
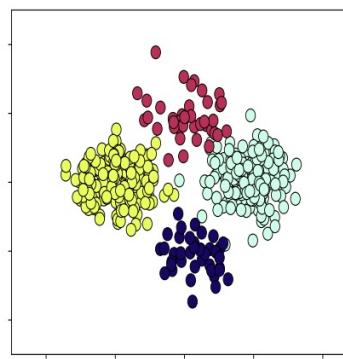


Self-supervised learning still suffers from data imbalance

- Performance degeneration: Linear probing on imbalanced data (D_i) and balanced data (D_b) with same data amount

Dataset	Subset	Many	Medium	Few	All
CIFAR10	D_b	77.14 ± 4.64	74.25 ± 6.54	71.47 ± 7.55	74.57 ± 0.65
	D_i	76.07 ± 3.88	67.97 ± 5.84	54.21 ± 10.24	67.08 ± 2.15
CIFAR100	D_b	25.48 ± 1.74	25.16 ± 3.07	24.01 ± 1.23	24.89 ± 0.99
	D_i	30.72 ± 2.01	21.93 ± 2.61	15.99 ± 1.51	22.96 ± 0.43

- Representation learning disparity: head classes dominate the feature regime but tail classes passively collapse



[1] Jiang et al. "Self-damaging contrastive learning." ICML 2021.

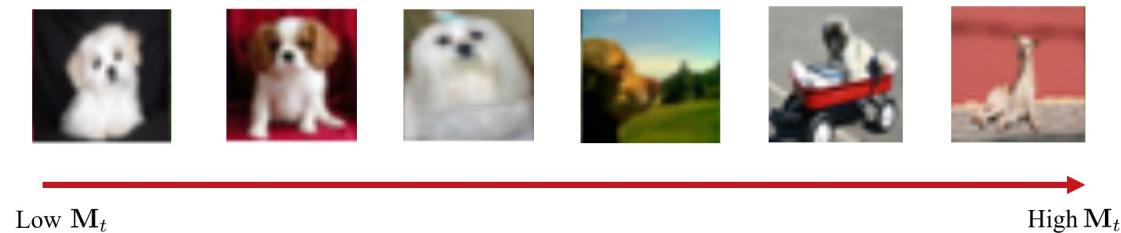
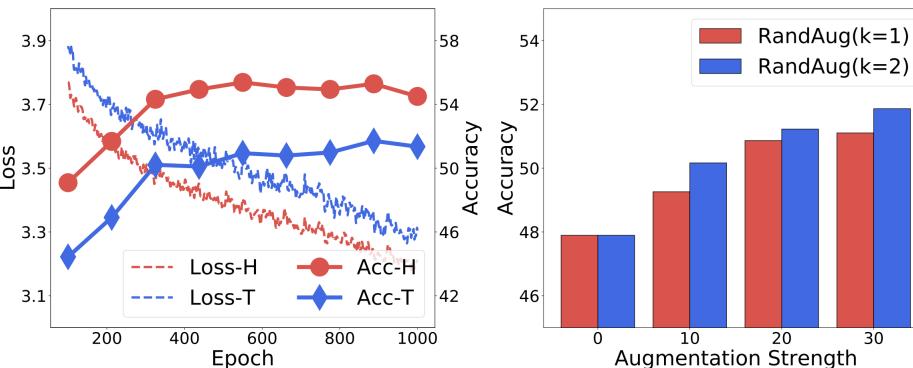
[2] Zhou et al. "Combating Representation Learning Disparity with Geometric Harmonization." NeurIPS 2023.



BCL: Boosted Contrastive Learning

- Motivation I: Memorization effect still holds under long-tailed distribution.
- Motivation II: Stronger information discrepancy motivates tail samples mining

➤ **Challenge:** how to detect tail data and how to construct the desired information discrepancy



- Motivated from the observation that *learning speed-based proxy* shows strong correlation with the memorization score[1], BCL extends the memorization estimation to *self-supervised learning*.

$$\mathcal{L}_{i,0}^m = \mathcal{L}_{i,0}, \quad \mathcal{L}_{i,t}^m = \beta \mathcal{L}_{i,t-1}^m + (1 - \beta) \mathcal{L}_{i,t} \quad \mathbf{M}_{i,t} = \frac{1}{2} \left(\frac{\mathcal{L}_{i,t}^m - \bar{\mathcal{L}}_t^m}{\max \{ |\mathcal{L}_{i,t}^m - \bar{\mathcal{L}}_t^m| \}_{i=0,\dots,N}} + 1 \right)$$

$$\Psi(x_i; \mathcal{A}, \mathbf{M}_i) = a_1(x_i) \circ \dots \circ a_k(x_i),$$

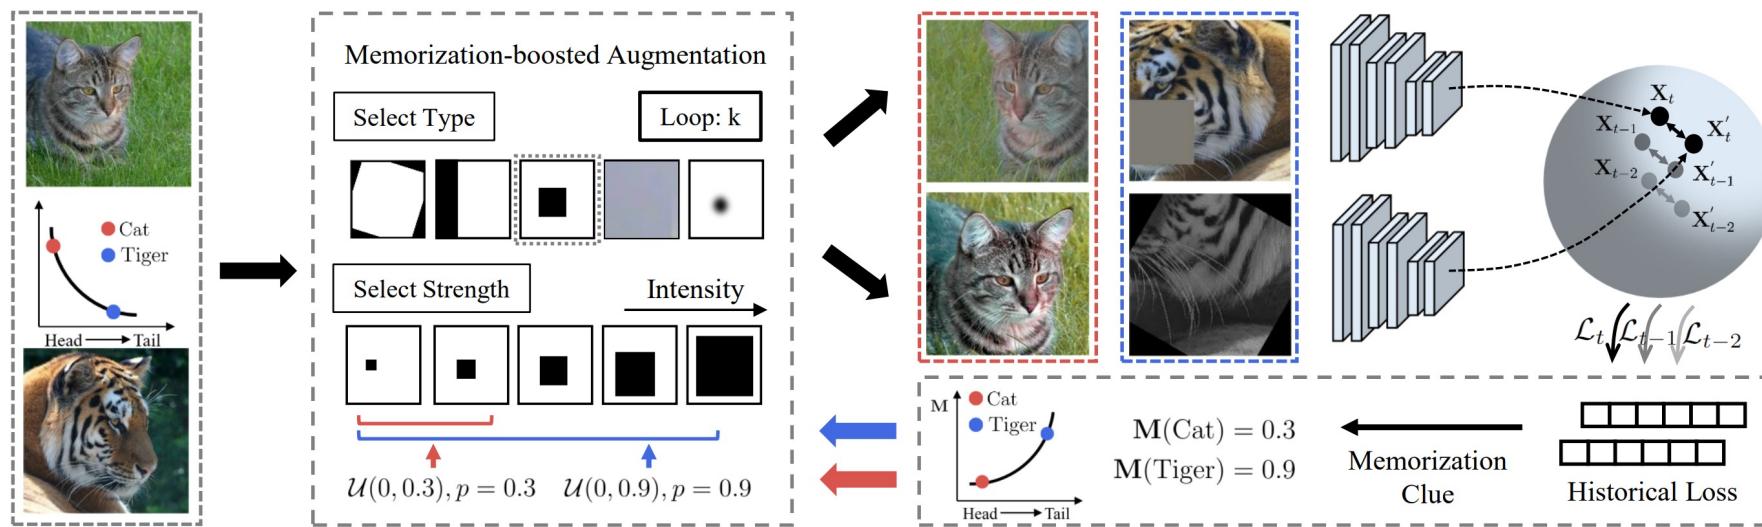
$$a_j(x_i) = \begin{cases} A_j(x_i; \mathbf{M}_i \zeta) & u \sim \mathcal{U}(0, 1) \text{ \& } u < \mathbf{M}_i \\ x_i & \text{otherwise} \end{cases} \quad \mathcal{L}_{\text{BCL}} = \frac{1}{N} \sum_{i=1}^N -\log \frac{\exp \left(\frac{f(\Psi(x_i))^T f(\Psi(x_i^+))}{\tau} \right)}{\sum_{x'_i \in X'} \exp \left(\frac{f(\Psi(x_i))^T f(\Psi(x'_i))}{\tau} \right)}$$

Adaptively assigns the appropriate augmentation strength for the individual sample according to the feedback from the memorization clues

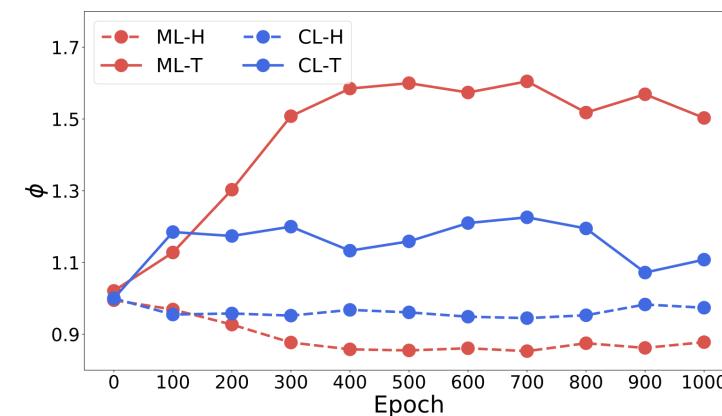
[1] Jiang et al. “Characterizing structural regularities of labeled data in overparameterized models.” ICML 2021 SJTU IL

[2] Zhou et al. “Contrastive learning with boosted memorization.” ICML 2022.

BCL: Boosted Contrastive Learning



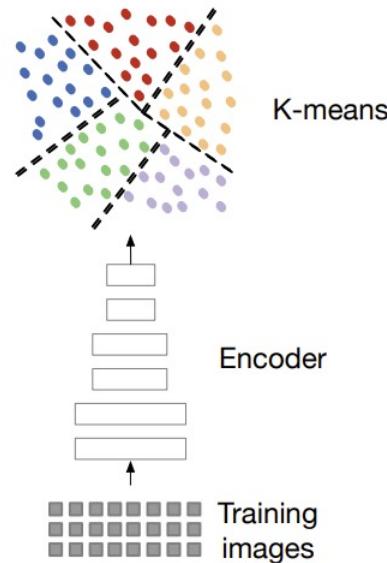
- Calculate **memorization scores** based on historical statistics to detect tail.
- Construct **instance-wise augmentations** to enhance representation learning.



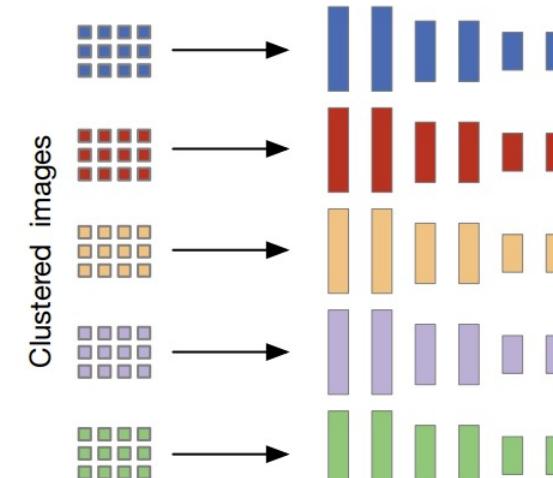
DnC: Divide and Contrast: Self-Supervised Learning from Uncurated Data

➤ **Motivation:** Conquer-and-Divide Training to isolate the negative effect on tail classes during training.

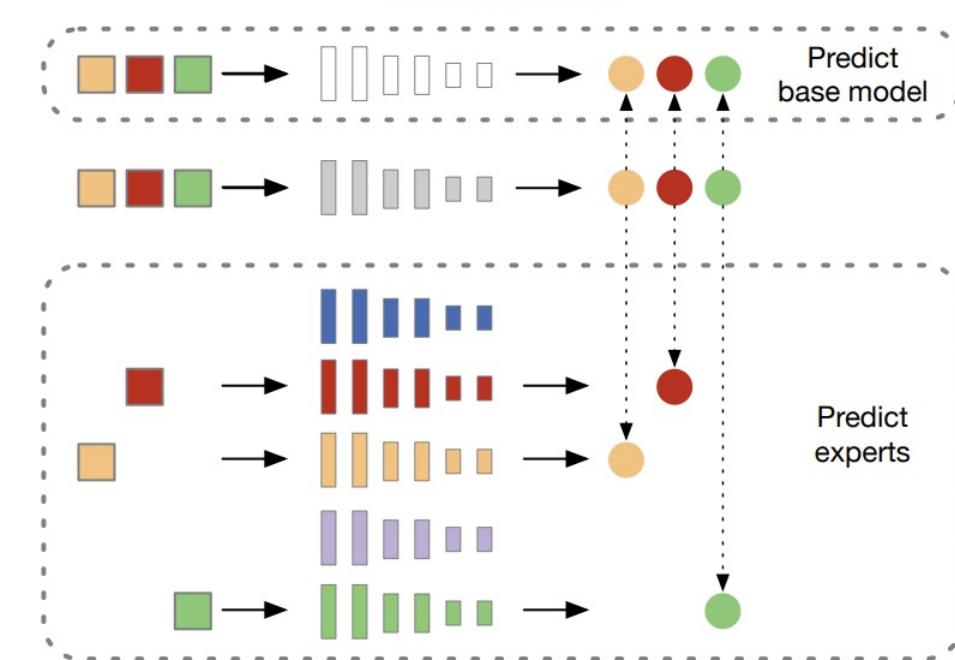
1. Train base model & cluster representations



2. Train expert models on subsets



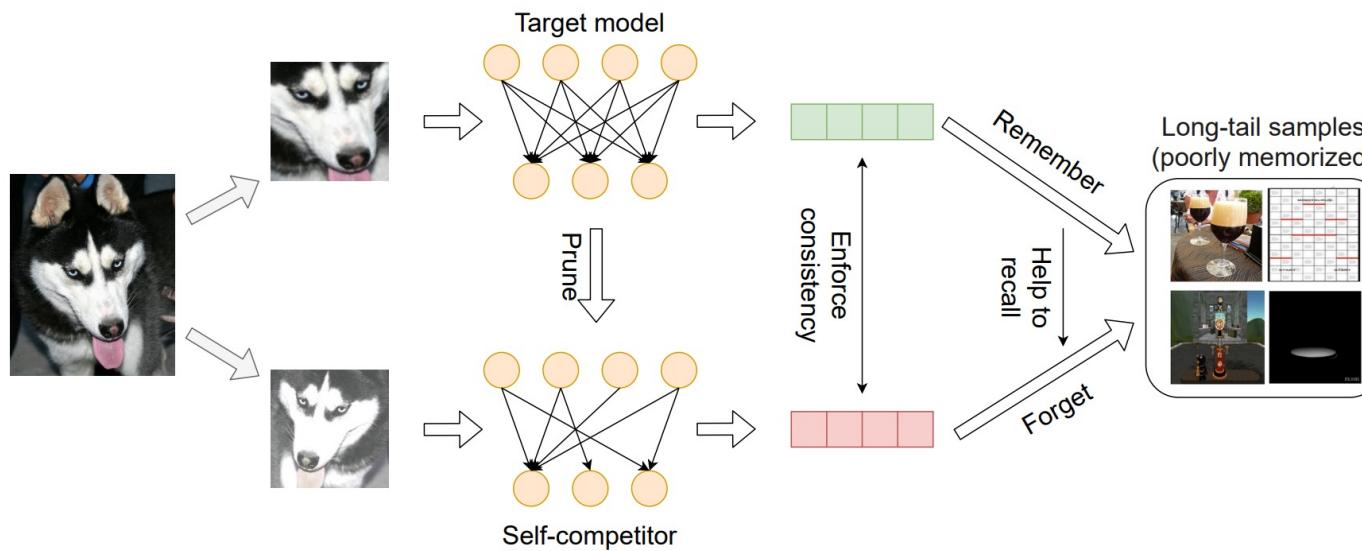
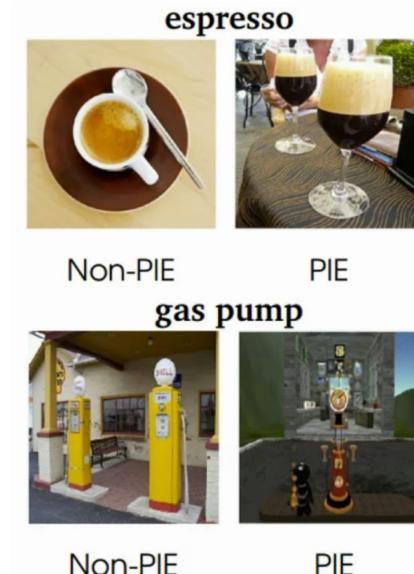
3. Distillation



SDCLR: Self-damaging Contrastive Learning

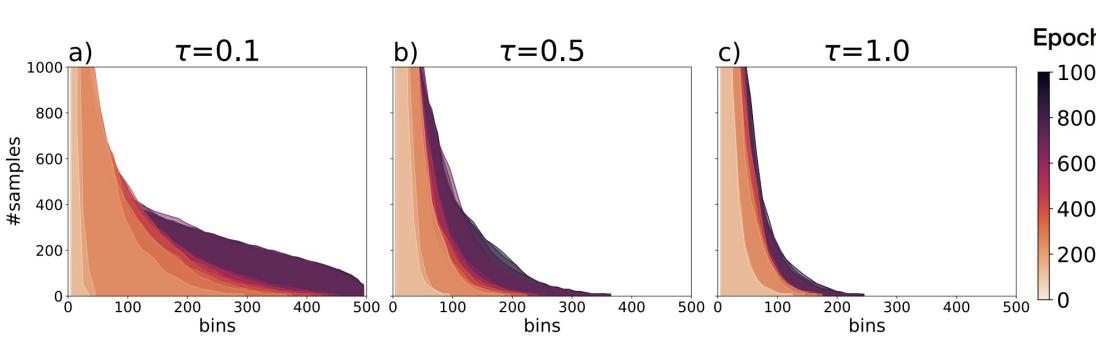
➤ **Intuition:** The sensitivity of head and tail samples to the model pruning, are very different, which helps us to anchor and promote the training of tail samples.

➤ **Pruning identified exemplars (PIE)** systematically investigates the model output changes introduced by pruning and finds that certain examples are particularly sensitive to sparsity. **They are highly likely to be rare and atypical samples, which probably comes from tail classes.**

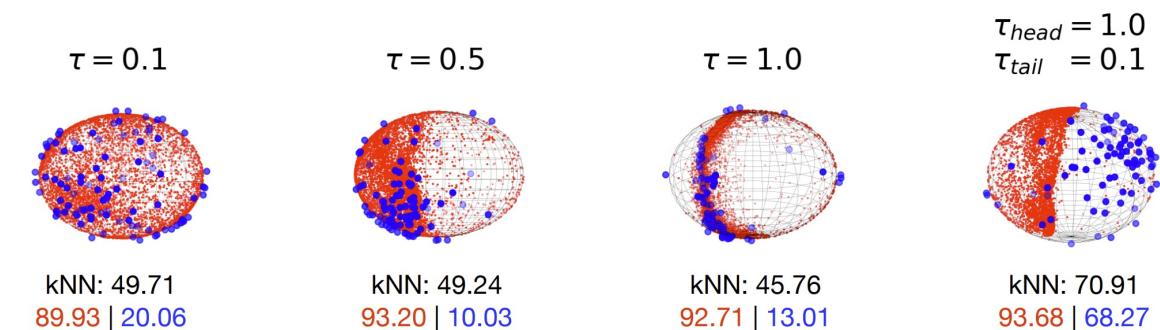


TS: Temperature Schedules for contrastive learning

- **Observation:** TS investigates the role of the temperature parameter τ in the contrastive loss, and find that a large τ emphasizes **group-wise discrimination**, whereas a small τ leads to a higher degree of **instance discrimination**.



Coverage of embedding space during training. For small τ the representations are more uniformly distributed.



Representations of a head and a tail class. **Red: head class** and **blue: tail class**. Small $\tau = 0.1$ promotes uniformity, while large $\tau = 1.0$ creates dense clusters. **Tail classes benefit from instance discrimination.**

- **Temperature Schedules (TS):** alternates between an upper τ and a lower τ bound at a fixed period length

$$\tau_{\cos}(t) = (\tau_+ - \tau_-) \times (1 + \cos(2\pi t/T))/2 + \tau_-$$

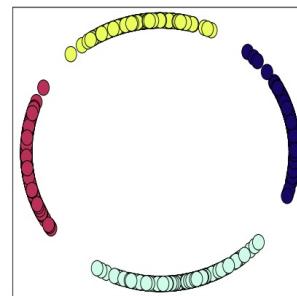
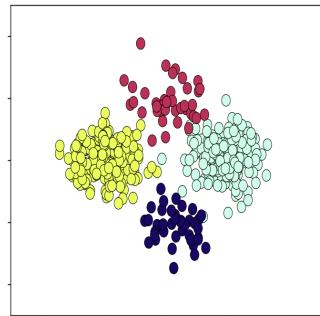
method	CIFAR-10-LT			CIFAR-100-LT			ImageNet-100-LT		
	kNN@10	FS LP	LS LP	kNN@10	FS LP	LT LP	kNN@10	FS LP	LT LP
SimCLR	60.19	68.29	61.68	28.12	25.70	31.20	38.00	42.64	44.82
SDCLR	60.74	71.03	64.99	29.22	27.28	34.23	37.36	42.74	46.40
SimCLR + TS	62.91	71.86	65.03	30.06	28.89	33.28	38.86	45.18	47.26

GH: Geometric Harmonization

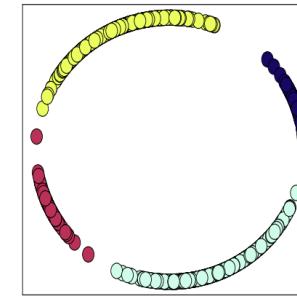
- Why the conventional contrastive learning underperforms in self-supervised long-tailed context?

Conventional contrastive loss motivates *sample-level uniformity*, which is biased towards the head classes.

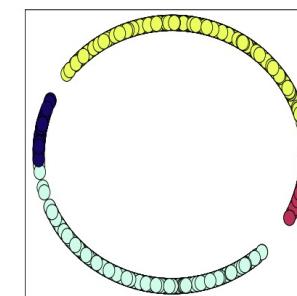
- Geometric Harmonization aims at achieves *category-level uniformity*, i.e., equal allocation w.r.t. classes



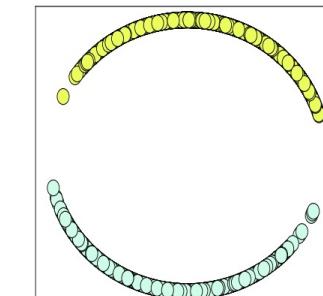
(a) R=1(Balanced)



(b) R=4



(c) R=16



(d) R=64

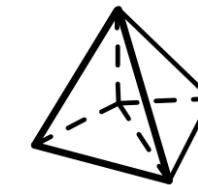
Contrastive learning causes severer representation learning disparity when enlarging the imbalance ratios.

- Challenges I: No guarantee for the desired category-level uniformity
- Challenges II: The latent true labels are not available, while the estimated labels are noisy

Geometric Uniform Structure

$$\mathbf{M}_i^\top \cdot \mathbf{M}_j = C, \quad \forall i, j \in \{1, 2, \dots, K\}, \quad i \neq j,$$

Any two vectors in \mathbf{M} have the same angle, namely, the unit space are equally partitioned by the vectors.



Surrogate Label Allocation

$$\min_{\hat{\mathbf{Q}}=[\hat{\mathbf{q}}_1, \dots, \hat{\mathbf{q}}_N]} \mathcal{L}_{\text{GH}} = -\frac{1}{|\mathcal{D}|} \sum_{\mathbf{x}_i \sim \mathcal{D}} \hat{\mathbf{q}}_i \log \mathbf{q}_i,$$

$$\text{s.t. } \hat{\mathbf{Q}} \cdot \mathbb{1}_N = N \cdot \pi, \quad \hat{\mathbf{Q}}^\top \cdot \mathbb{1}_K = \mathbb{1}_N,$$

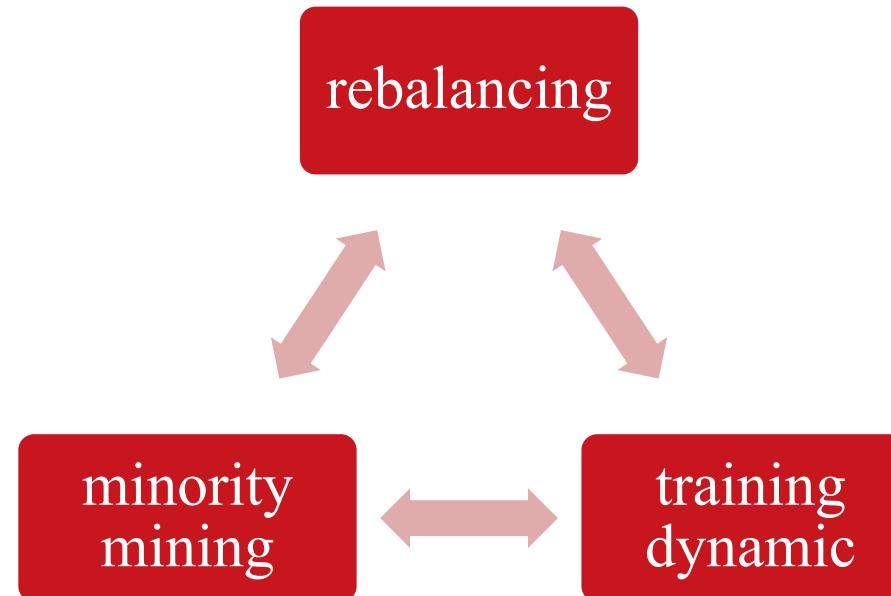
Overall objective

$$\min_{\theta, \hat{\mathbf{Q}}} \mathcal{L} = \mathcal{L}_{\text{InfoNCE}} + w_{\text{GH}} \mathcal{L}_{\text{GH}},$$

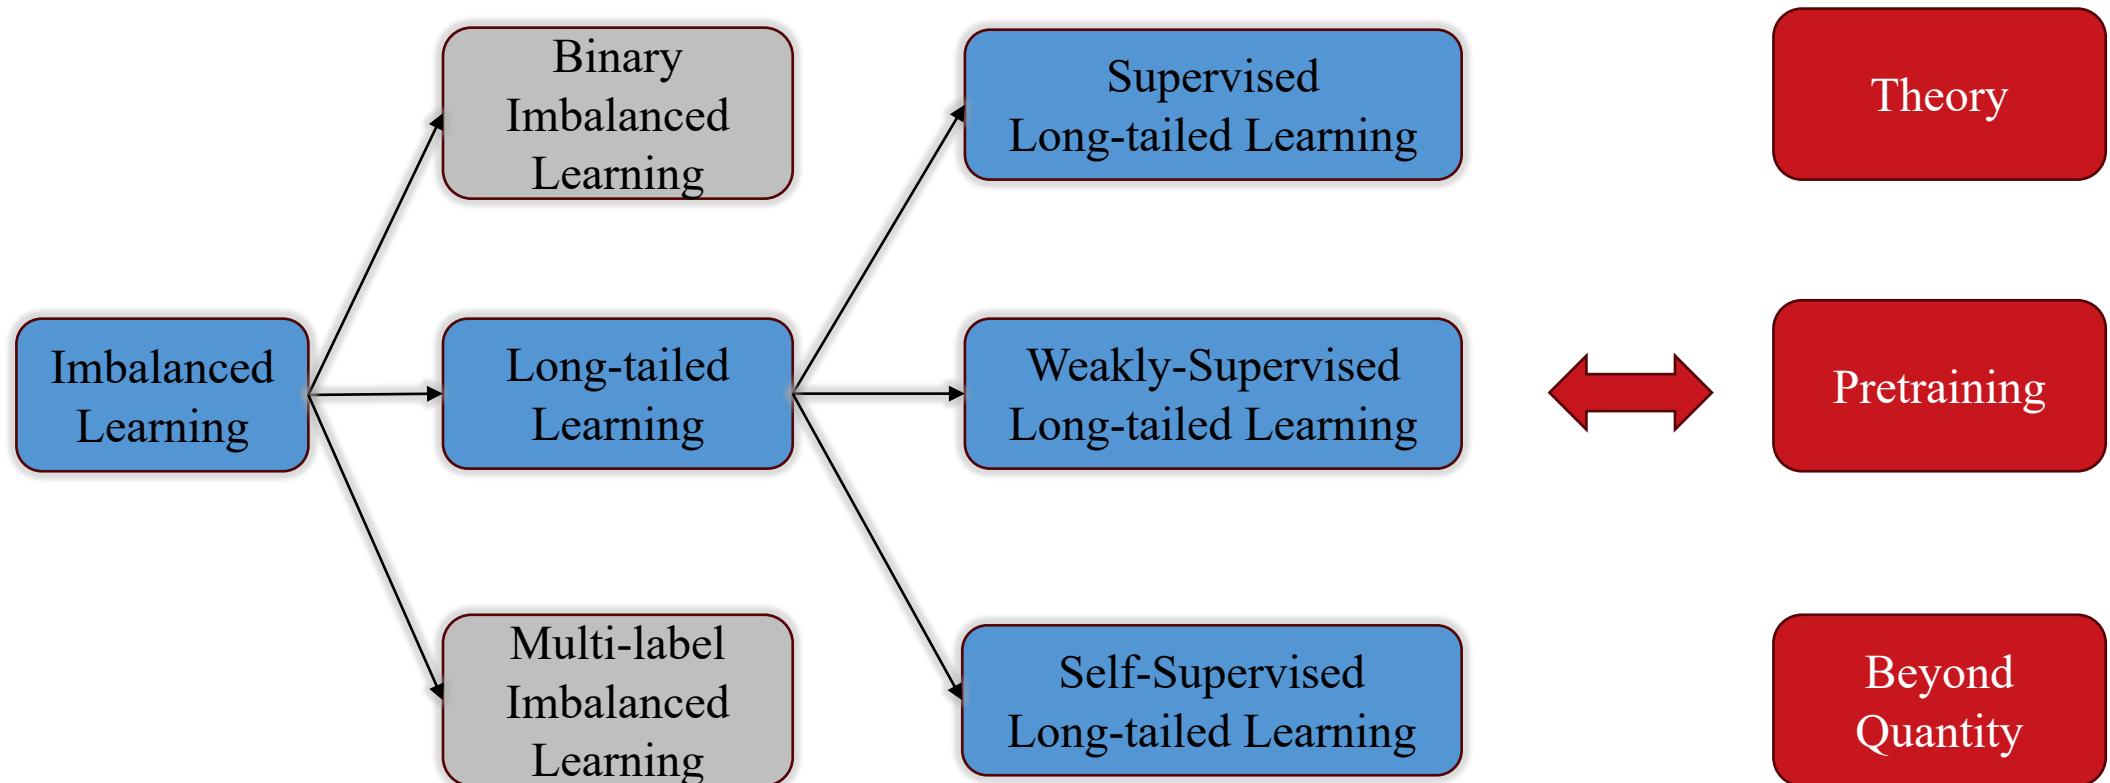


Summary

- Self-supervised learning is **more robust** to data imbalance **than the supervised counterpart**
- However, self-supervised learning still suffers from the long-tailed distribution, resulting in performance degeneration and **representation learning disparity**



Still require more efforts on this way





Thank you

Q & A

Tutorial website: <https://tmlr-group.github.io/tutorials/aaai2024.html>

